

UNIVERSITY OF UDINE

DEPARTMENT OF MATHEMATICS, COMPUTER SCIENCE AND PHYSICS

READERSOURCING 2.0:
DOCUMENTATION
VERSION 2.0.0

MICHAEL SOPRANO AND STEFANO MIZZARO

Contents

List of Figures	iii
List of Tables	iv
1 Introduction	1
2 General Architecture	1
3 RS_Server	1
3.1 Implementation and Technology	2
3.2 Communication Paradigm	3
3.3 Database	6
3.4 Class Diagram	7
3.5 Deploy	13
3.5.1 Modality 1: Manual	13
3.5.1.1 Requirements	13
3.5.1.2 How To	13
3.5.1.3 Quick Cheatsheet	14
3.5.2 Modality 2: Docker Compose	14
3.5.2.1 Requirements	14
3.5.2.2 How To	14
3.5.3 Environment Variables	15
3.5.4 Setting Variables	17
3.5.5 Sending Mails	18
3.5.6 Connecting To The Database	19
3.5.7 Logging To The Standard Output	19
4 RS_PDF	19
4.1 Implementation and Technology	19
4.2 Package Diagram	20
4.3 Class Diagram	21
4.4 Installation	21
4.4.1 Requirements	21
4.4.2 Command Line Interface	22
5 RS_Rate	23
5.1 Implementation and Technology	23
5.2 Installation	23
5.3 Development and Packaging	24
5.4 Usage	24

6	RS_Py	29
6.1	Implementation and Technology	29
6.2	Installation	30
6.3	Usage	30
7	Overview	31
	References	33

List of Figures

1	Conceptual architecture of Readersourcing 2.0 (NOT UML). RS_PDF is bundled with and invoked by RS_Server.	2
2	Intuitive scheme of the MVC pattern (NOT UML).	3
3	Persistent entities and associations in the RS_Server database (selected attributes).	7
4	General class diagram of RS_Server. Method lists are intentionally selective. .	8
5	Authentication process of the web interface. REST clients use the returned JWT through the Authorization header.	12
6	Package diagram of RS_PDF.	20
7	General class diagram of RS_PDF v2.0.0.	21
8	RS_Rate characterized as an extension having a popup action.	25
9	The login page of RS_Rate.	25
10	The rating page of RS_Rate.	26
11	The user registration page of RS_Rate.	27
12	The rating page of RS_Rate after a “save for later” request.	28
13	A publication annotated through RS_Rate.	28
14	The server interface to rate a publication.	29
15	The profile page of RS_Rate.	30
16	Readers’ interaction modalities with the Readersourcing 2.0 ecosystem. . . .	32

List of Tables

1	Subset of the RESTful interface of RS_Server.	5
2	Environment variables of RS_Server.	17
3	Command line options of RS_PDF.	22

1 Introduction

The Readersourcing 2.0 ecosystem was developed as part of a research project funded jointly by SISSA Medialab¹ and the University of Udine.² It was presented by Soprano et al. [5] at the IRCDL 2019 Conference.³ The paper is also available on Zenodo.⁴ The code and related documentation are available on GitHub.⁵

Initially, a recap of its general architecture is presented, followed by a brief description of the role and purpose of each of its components. Specific aspects, such as the technology used, the internal architecture, and the structure of the database, are also discussed. This is achieved through different types of diagrams adhering to the UML standard unless otherwise specified, drawn according to the style rules proposed by Fowler [2]. The current software requirements and operational procedures are recorded alongside the original architectural description.

2 General Architecture

Readersourcing 2.0 is an ecosystem composed of several components. It includes RS_Server [9] (presented in Section 3), which gathers the ratings given by readers, and RS_Rate [8] (presented in Section 5), which acts as a client and enables readers to rate publications. Although all operations can be carried out directly through the web interface provided by RS_Server, RS_PDF [6] (presented in Section 4) annotates PDF files with a rating URL and QR code on behalf of the server application. Finally, RS_Py (presented in Section 6) provides a fully functional Python implementation of the RSM and TRM models described by Soprano et al. [5], together with the notebooks used to generate datasets and analyze experiments.

Figure 1 provides an overview of the architecture of Readersourcing 2.0, and in the following, we briefly describe these four components. We summarize the capabilities of the ecosystem in Section 7.

3 RS_Server

RS_Server [9] is the server component responsible for collecting and aggregating the ratings given by readers and applying the RSM and TRM models described by Soprano et al. [5] to compute quality scores for readers and publications. An instance of RS_Server is deployed with a compatible RS_PDF executable, which is shipped in the `lib` directory and invoked as a local process when a publication must be annotated. Up to n browsers and

¹<https://medialab.sissa.it/>

²<https://www.uniud.it/>

³<https://ircdl2019.isti.cnr.it/>

⁴<https://zenodo.org/record/1446468>

⁵<https://github.com/Miccighel/Readersourcing-2.0>




Figure 1: Conceptual architecture of Readersourcing 2.0 (NOT UML). RS_PDF is bundled with and invoked by RS_Server.

their users can communicate with the server, each through an instance of RS_Rate. The extension is distributed for browsers based on Chromium and for Firefox.

Readers can interact with the server either through clients installed in their browsers or through the standalone web interface provided by RS_Server. These clients manage reader registration and authentication, rating actions, and downloads of annotated publications.

During the design phase of RS_Server, strategies were adopted to ensure extensibility and generality. This includes:

- (i) Straightforward addition of new models.
- (ii) A shared input data format among all models.
- (iii) A standard procedure for models needing to save values locally in the RS_Server database.

3.1 Implementation and Technology

RS_Server is developed using *Ruby on Rails*,⁶ an open source web application framework built on the Ruby programming language and strongly based on the Model View Controller (MVC) architectural pattern.

The MVC pattern separates a program’s control logic from its data presentation and business logic, helping developers establish a clear architecture from the initial design phase.

⁶<https://rubyonrails.org/>




Figure 2: Intuitive scheme of the MVC pattern (NOT UML).

Figure 2 illustrates the resulting structure. It comprises three distinct entities: the *Controller*, which manages control logic; the *Model*, which encapsulates business logic; and the *View*, which presents data.

The Controller has direct access to both the Model and the View. It typically receives user input from the View, updates the internal state of the Model through its methods, and then passes the updated Model back to the View for presentation. A software system may contain several Controllers, each of which can manage multiple Model instances. In web MVC frameworks such as Rails, Controllers are commonly organized around domain resources, and a Model can be presented through more than one View.

MVC is not the only principle underlying Rails. Another central principle is “Convention Over Configuration”: the framework reduces the number of decisions required during application development by providing standard conventions that can still be changed when greater flexibility is needed. The “Rails Doctrine” describes these principles in greater detail.⁷

Rails is an actively developed framework used by organizations such as *GitHub* and *SoundCloud*. It is supported by an established community and extensive learning resources.

3.2 Communication Paradigm

An MVC framework such as Rails can support several kinds of web applications, including a *Web Service*. A Web Service exposes operations remotely by exchanging messages encoded in a standard format such as *JSON* over a transport protocol such as *HTTP*. Its interface must define the available resources and operations, together with the messages required to access them.

⁷<https://rubyonrails.org/doctrine/>

One communication paradigm for Web Services is *REST* (*REpresentational State Transfer*). In a RESTful interface, resources are identified by URIs and the HTTP method determines the operation to perform. The server returns a response in the agreed interchange format, and the client interprets that response.

RS_Server is a Rails application exposing RESTful interfaces alongside web workflows rendered by the server. API clients exchange JSON messages over HTTP, while the same domain and authorization rules are reused by the web interface.

The REST interface is defined by the application routes, JSON views, and integration tests. A Postman collection provides request and response examples for the exposed resources:

<https://documenter.getpostman.com/view/4632696/2sBYAswWvk>.

For example, Table 1 presents a subset of the RESTful interface of RS_Server. These operations manage publications, one of the entities in the application domain. If a user requests the “show” operation for publication 1 through the corresponding endpoint, RS_Server returns a JSON response similar to the following example.

```
1 {
2   "id": 1,
3   "doi": "10.1140/epjc/s10052-018-6047-y",
4   "title": "Uncertainties in WIMP dark matter scattering revisited",
5   "subject": null,
6   "author": "John Ellis",
7   "creator": "Springer",
8   "producer": null,
9   "pdf_url": "https://example.org/publication.pdf",
10  "steadiness": 1.0,
11  "score_rsm": 0.71,
12  "score_trm": 0.68,
13  "created_at": "2026-01-01T12:00:00.000Z",
14  "updated_at": "2026-01-01T12:00:00.000Z",
15  "url": "http://localhost:3000/publications/1.json"
16 }
```

Publication records are shared among readers. The REST interface therefore exposes their creation and retrieval, together with dedicated domain operations such as fetching and refreshing, but does not expose generic update or deletion operations. This prevents one authenticated reader from arbitrarily rewriting or removing a shared publication and its associated ratings.

Operations that regenerate annotated files or change a reader’s subscription preference use POST. An unsubscribe link opens a confirmation page, and the preference changes only after the authenticated reader submits that page.

Endpoint	HTTP Message	Operation	Description
/publications.json	GET	Index	Fetches the entire collection of Publications.
/publications/list	GET	List	Fetches the entire collection of Publications (view rendered by the server).
/publications/1.json	GET	Show	Returns the Publication with identifier equal to 1.
/publications/lookup.json	POST	Lookup	Searches for a Publication by its URL without fetching it.
/publications/random.json	GET	Random	Returns a random Publication.
/publications/1/is_rated.json	GET	Is Rated	Checks if the Publication with identifier equal to 1 has been rated by at least one reader.
/publications/1/is_saved_for_later.json	GET	Is Saved For Later	Checks if the Publication with identifier equal to 1 has been saved for later by the current user.
/publications.json	POST	Create	Creates a new Publication.
/publications/is_fetchable.json	POST	Is Fetchable	Retrieves and opens the bounded source as a PDF for API inspection.
/publications/extract.json	POST	Extract	Extracts the rating URL from an annotated Publication.
/publications/fetch.json	POST	Fetch	Retrieves, annotates, verifies, and publishes a Publication.
/publications/fetch_upload.json	POST	Fetch Upload	Performs the same preparation from an uploaded original PDF.
/publications/1/refresh.json	POST	Refresh	Fetches the Publication with identifier equal to 1 again.
...	...	...	...

Table 1: Subset of the RESTful interface of RS_Server.

3.3 Database

To support features such as storing annotated publications and authenticating users, RS_Server uses the database schema illustrated in Figure 3. The Readersourcing 2.0 application domain contains three main entities:

- **Users:** represents the users of the system and stores their personal data. A user may have an optional *ORCID* and a boolean indicating whether they wish to receive emails. Additional attributes store the tokens required for operations such as password reset.
- **Ratings:** represents the ratings provided by readers for publications. Each rating has an integer score from 0 to 100, including both endpoints, and an original score that retains the first submitted value when the rating is edited. Two booleans indicate whether the rating is anonymous and whether it was edited later. A reader can provide at most one rating for a given publication; this invariant is enforced both by the application model and by a unique database constraint.
- **Publications:** represents publications known to the server, including those saved for later but not yet rated. A publication may have a *DOI*, descriptive metadata, and attributes used to manage the original and annotated PDF files.

These entities also contain attributes such as *steadiness* and *informativeness*, which store scores or parameters computed by the Readersourcing models.

Both rating scores are required and must remain within the rating scale. The application validates the supplied scores, and the database enforces their presence and range. This applies to ratings submitted through the API or an annotated PDF and to edits of existing ratings. An invalid value leaves the stored domain scores unchanged and does not send a confirmation email; JSON requests return HTTP status 422 with validation messages for the affected fields. Before installing the database constraints, the migration checks existing ratings and stops if their scores are missing or out of range, without correcting those records automatically.

An auxiliary **AuthenticationToken** entity records each active login. It belongs to one user and stores the unique JWT identifier and its expiration time, but never the encoded JWT itself. Authentication tokens support the security boundary rather than the Readersourcing score domain.

Figure 3 depicts the two mandatory associations of a rating. A user may give zero or more ratings, but each rating belongs to exactly one user. An anonymous rating still retains this association so that authentication and the constraint of one rating per user can be enforced; anonymity controls how the reader identity participates in presentation and processing, not referential integrity.

Similarly, a publication may receive zero or more ratings, while each rating belongs to exactly one publication. The zero multiplicity is required because a publication may be fetched and saved for later before any rating is submitted.




Figure 3: Persistent entities and associations in the RS_Server database (selected attributes).

The auxiliary association permits a user to have zero or more active authentication tokens, while every token belongs to exactly one user. Removing a user also removes all their active tokens.

3.4 Class Diagram

Figure 4 shows the main responsibilities of the RS_Server implementation. Rails controllers map HTTP operations to the domain models, while JSON and views rendered by the server present their results. The three persistent domain models encapsulate the Readersourcing logic and associations described in the previous section. The auxiliary *AuthenticationToken* model maintains the registry of active logins, while objects that are not persisted represent durable paper references and request rate policies. Authentication, reference validation, request throttling, and Readersourcing computations are delegated to focused objects rather than being absorbed into the controllers.

The *ApplicationController* contains shared authorization hooks for both API and server requests, while the *AuthenticationController* handles login and logout. The *Authenticator*, *Authorizer*, and *AuthenticationTokenRevoker* command objects isolate credential verification, token issuance, token validation, and revocation from HTTP concerns.

The *PaperRatingReference* value object similarly isolates the creation and validation of references stored in annotated publications. It uses authenticated encryption with a dedicated purpose and binds each reference to one reader and one publication without persisting a login credential in the PDF.

The *PublicationDownloadReference* value object authorizes access to a prepared copy without exposing its private storage path. A signed reference is bound to one reader,




Figure 4: General class diagram of RS_Server. Method lists are intentionally selective.

publication, PDF variant, and file name, and expires after five minutes by default. This separate capability allows both browser interfaces to open a PDF in a new tab, where an **Authorization** header cannot be added, without making the storage directory public.

The *RequestRateLimit* value object describes each request budget and provides opaque identities for the corresponding cache counters. Controllers apply these policies through the Rails rate limiting boundary, leaving credential verification, mail delivery, publication handling, and the Readersourcing domain objects unaware of throttling concerns.

For REST clients, relying only on session data held by the server to identify the authenticated user is unsuitable because the RESTful interface is *stateless*. Clients therefore attach an authentication *token* to each protected request. The workflow rendered by the server can still keep the same token in its Rails session.

Authentication establishes the identity of the reader making the request, while authorization at the resource level determines which records that reader may operate on. Profile updates, account deletion, and subscription changes are restricted to the current reader. A reader may inspect and edit only their own rating through the individual rating endpoints. Reader collections and individual reader responses expose the public academic profile—name, ORCID, and Readersourcing scores—without email or subscription preferences; the current reader obtains those private profile fields through the dedicated information endpoint.

When a user logs in, the supplied credentials are verified against the stored password digest and the account confirmation state. The credentials themselves are not placed in the token. If verification succeeds, RS.Server creates an active token record and issues a JWT signed with HS256 whose payload contains the user identifier, the client IP address, its unique *jti* identifier, and a standard expiration claim set to seven days after issuance.

The authentication response returns the JWT for REST clients. These clients attach it to protected requests through the **Authorization** header, which accepts both the direct token format used by RS.Rate and the conventional **Bearer <token>** format. The web interface instead keeps the token only in the encrypted Rails session, which is implemented through a cookie marked **HttpOnly**; browser scripts neither read the JWT nor add it to request headers. Upon receiving a protected request, the *Authorizer* verifies the JWT signature, expiration time, IP address, referenced user, and corresponding active token record before allowing the request to proceed. Figure 5 provides an intuitive illustration of this authentication process.

The Rails session identifier is renewed after successful authentication, before the token is stored, and is renewed again when the reader logs out. Logout removes the active token record for the current browser or API session and leaves the reader’s other devices authenticated. A successful password change or reset removes every active token belonging to that reader; account deletion removes the same records through the user association. The signed value remains immutable, but without its active record it can no longer authorize a protected request.

Logout changes application state and is therefore available only through an HTTP POST request. HTML requests that change state require a valid Rails authenticity token. JSON

requests from the web interface use the Rails session and require the same authenticity token. JSON requests that do not rely on that session remain stateless; protected API requests carry the JWT through **Authorization** and are handled through the token authentication boundary.

In the save for later workflow, RS_Server places a dedicated paper rating reference in the URL supplied to RS_PDF. Its encrypted payload contains the reader and publication identifiers and has no automatic expiration. It is separate from the JWT, which is valid for seven days, and cannot authorize an API request. Following the reference always requires a current authenticated session, after which the server verifies that the reference belongs to both the selected publication and the authenticated reader. Logout, password changes, and password resets therefore do not invalidate an annotated PDF. Removing its reader or publication makes the reference unusable.

The web interface and RS_Rate do not download a publication merely because a reader enters a URL or opens the extension. They first report that the source is ready to prepare, and the real operation starts when the reader requests the prepared PDF. RS_Server then retrieves the source once under the configured address, redirect, timeout, and size policies. The declared content type is retained as an indication, while the file header and the PDF parser determine whether the received content is a usable PDF. Private and reserved network addresses remain refused unless a controlled deployment enables the documented override.

The same preparation object is used for remote and uploaded sources. An upload provides an alternative when a publisher requires access through the reader's browser session or when the remote source is otherwise unavailable. The source is opened before RS_PDF is invoked and a PDF that is malformed, encrypted, unsupported, or already contains a **BaseUrl** rating reference is refused with a specific status. RS_PDF writes into temporary storage. Before the generated file replaces any existing reader copy, RS_Server opens it again and verifies that it contains exactly one additional page and the expected rating URL. A failed process or final verification leaves the previous copy untouched. Completed copies remain outside the public asset directory and are served only after a download reference has been verified. The response forbids shared caching and prevents the reference from being sent as a referrer. Both user interfaces translate the stable preparation statuses into an explicit title, explanation, retry action, and upload option where applicable.

To recover a rating link from an annotated publication, the reader uploads the PDF to the extraction endpoint. The upload follows the configured size limit and is opened as a PDF before its **BaseUrl** metadata is read. The address must use the exact public origin and rating route of RS_Server, without credentials, a query, or a fragment. Its encrypted reference must identify an existing publication and the authenticated reader. Invalid files or references return HTTP status 422, and the temporary copy is removed when the request finishes. The web interface also verifies the origin and route before opening the returned link and displays an error if validation fails.

Password recovery uses a separate credential that can be used once. An HTTP **GET** request only displays the recovery form; sending the request requires **POST**. The request

endpoint returns the same confirmation for registered and unregistered addresses. For a registered address, RS_Server generates a random token, stores only its SHA-256 digest and issue time, and sends a link built from the configured public origin. The link remains valid for four hours. It opens a form rendered by the server in which the reader chooses and confirms a new password; a successful reset clears the stored digest before a confirmation email is sent. Passwords are never sent by email.

Request rate policies bound operations that can be abused or consume disproportionate resources. By default, one IP address may make ten authentication attempts in three minutes and five password recovery requests in fifteen minutes; one normalized email address may receive three recovery requests in thirty minutes; one IP address may send five contact messages in ten minutes; and one authenticated reader may initiate thirty combined publication download, inspection, extraction, and annotation operations in one hour. Exceeding a budget returns HTTP status 429 `Too Many Requests` with a `Retry-After` header and does not execute the protected action. Cache identities are SHA-256 digests rather than plain email addresses or IP values. Displaying the password recovery form does not consume a recovery request.

Every HTML and JSON response carries an enforced Content Security Policy. Browser scripts, styles, and fonts are installed through Yarn and served by the Rails asset pipeline, so executable code and presentation resources do not depend on external hosts at request time. Direct versions are declared in `package.json` and the complete dependency graph is recorded in `yarn.lock`. Visual dependencies remain within the compatibility lines used by the original interface. DataTables, JSZip, and js-cookie use the first versions that address their known registry advisories while preserving the APIs used by RS_Server. Pdfmake uses the 0.2 compatibility line, which preserves the API required by the original DataTables integration. Pdfmake and its font data are loaded only on the authenticated publication and reader lists, where they provide the original PDF export behaviour without allowing dynamic code evaluation. The original views and visual assets require style attributes, which remain permitted. Images are restricted to the application, data URLs, and the badge host used on the resources page. Other directives forbid object content, external form targets, external base URLs, and framing.

Responses set `X-Frame-Options` to `DENY`, `X-Content-Type-Options` to `nosniff`, and `Referrer-Policy` to `same-origin`. The `Permissions-Policy` header disables access to the camera, screen capture, location, microphone, payment, and USB interfaces. Rails adds HTTP Strict Transport Security only when `FORCE_SSL=true`.




Figure 5: Authentication process of the web interface. REST clients use the returned JWT through the `Authorization` header.

The classes responsible for the Readersourcing computations follow the *Strategy* pattern: *Readersourcing*, *ReadersourcingStrategy*, *RsmStrategy*, and *TrmStrategy*. A rating invokes the RSM strategy first and the TRM strategy second. The pattern keeps model selection extensible while expressing the score computation order explicitly.

The *RatingSubmission* command object coordinates the persistence of ratings and the results produced by these strategies. It saves a new rating and the RSM and TRM results in one database transaction. Rating submissions and edits acquire the same PostgreSQL advisory transaction lock: reader reputation is shared across publications, so locking only the selected publication would not protect every affected score. The reader and publication are read again after the lock is acquired. Each strategy requires its writes to succeed; a failed calculation rolls back the rating, any publication created for it, and all associated score changes. The lock coordinates rating operations across application processes and is released when their transaction ends. Editing a rating retains its original score and does not recompute either model.

Confirmation mail is attempted for subscribed readers only after the new rating and its calculations have been committed. A mail delivery failure leaves the saved rating intact and does not change the successful API response. The paper rating interface reports that the rating was saved but its confirmation email could not be sent.

3.5 Deploy

RS.Server supports two setup paths: a manual installation, intended primarily for development, and a Docker Compose deployment, which provides a reproducible application and database environment. Docker Compose can use the published `miccighel/rs_server:v2.0.0` image available on Docker Hub,⁸ or build the application locally from the same source tree.

Secrets and database settings specific to each environment are described in Section 3.5.3. A local `.env` file must never be committed.

3.5.1 Modality 1: Manual

This modality downloads and initializes RS.Server directly on the local machine and offers the greatest flexibility for development and domain model experiments.

3.5.1.1 Requirements

- Ruby == 3.4.10;
- Bundler compatible with the committed `Gemfile.lock`;
- a Java 21 runtime, required by RS.PDF;
- PostgreSQL >= 17;
- Node.js >= 18, used to install and precompile the existing browser assets.

3.5.1.2 How To

Clone the RS.Server repository,⁹ navigate to its root directory, copy `.env.example` to `.env` as described in Section 3.5.4, and run:

```
bundle install
node .yarn/releases/yarn-3.6.3.cjs install --immutable
bin/rails db:create
bin/rails db:migrate
bin/rails db:seed # optional sample data
bin/rails server -b 127.0.0.1 -p 3000 -e development
```

PostgreSQL must be accepting connections before the database commands are executed. With the sample bind address and port, the application is available at `http://127.0.0.1:3000`. Rails binstubs are always invoked from the repository root so that the locked application dependencies are used.

The active token registry is populated at login. A session carrying a JWT without a registered identifier is rejected and the reader signs in again.

⁸https://hub.docker.com/r/miccighel/rs_server

⁹https://github.com/Miccighel/Readersourcing-2.0-RS_Server

3.5.1.3 Quick Cheatsheet

1. change to the main directory with `cd`;
2. copy `.env.example` to `.env` and populate it;
3. `bundle install`;
4. `node .yarn/releases/yarn-3.6.3.cjs install --immutable`;
5. `bin/rails db:create`;
6. `bin/rails db:migrate`;
7. `bin/rails db:seed` (optional);
8. `bin/rails server -b <address> -p <port> -e development`.

3.5.2 Modality 2: Docker Compose

Docker Compose uses the public `miccighel/rs_server:v2.0.0` image and starts it together with PostgreSQL 17. The application image includes Ruby 3.4.10, the Java 21 runtime required by RS.PDF, the locked Ruby dependencies, and the precompiled browser assets. When the source code has been edited, the same image can be built locally. This is the recommended reproducible deployment for local integration and verification under conditions similar to production. Named volumes retain the database and the prepared publications when the containers are recreated. The application runs as an unprivileged user with a read only root file system; temporary directories and the private publication volume are its only writable locations.

3.5.2.1 Requirements

- Docker Desktop or another installation supporting Docker Compose v2.¹⁰

3.5.2.2 How To

Clone the repository, copy `.env.example` to `.env` and replace its example values, ensure that the Docker Engine is running, and execute:

```
docker compose up
```

The command downloads the published image when it is not already available. To build the image from the checked out source instead, execute `docker compose up --build`.

The `database` service must pass its health check before the `rs_server_setup` service runs `bin/rails db:migrate`. The setup service must complete successfully before

¹⁰<https://docs.docker.com/compose/>

`rs_server_webapp` starts; database migration is not performed by the application entry-point and sample data are not created implicitly. Once startup completes, the web application is available at `http://localhost:3000`; `0.0.0.0` is the container bind address, not the address users should enter in a browser. The `GET /up` endpoint can be used to verify that Rails boots successfully.

Sample data remain optional:

```
docker compose run --rm rs_server_webapp bin/rails db:seed
```

Use `docker compose down` to stop the services. The named volumes are retained so that database records and prepared publications survive ordinary container recreation.

Quick Cheatsheet

- change to the main directory with `cd`;
- copy `.env.example` to `.env` and populate it;
- `docker compose up`;
- `docker compose up --build` when the source code has been edited;
- `docker compose run --rm rs_server_webapp bin/rails db:seed` (optional);
- `docker compose down` (to stop the services).

3.5.3 Environment Variables

Credentials and connection details specific to each deployment are supplied through environment variables. Table 2 lists the variables used by the application. Database URLs take precedence over individual PostgreSQL settings.

Environment Variable	Purpose	Scope
<code>SECRET_DEV_KEY</code>	Rails secret used in development.	development
<code>SECRET_PROD_KEY</code>	Rails secret used to sign and encrypt production data, including JWTs, session data, and paper rating references. Keep it stable while issued references must remain usable.	production
<code>DATABASE_URL</code>	Complete PostgreSQL connection URL. Overrides the individual <code>POSTGRES_*</code> variables.	development, production
<code>TEST_DATABASE_URL</code>	Complete PostgreSQL connection URL for the test database.	test

Environment Variable	Purpose	Scope
POSTGRES_USER	PostgreSQL user name used when a full URL is not supplied.	all database environments
POSTGRES_PASSWORD	Password for the PostgreSQL user.	all database environments
POSTGRES_HOST	PostgreSQL host; defaults to <code>localhost</code> outside Compose and is set to <code>database</code> by Compose.	all database environments
POSTGRES_DB	Development or production database name; defaults to <code>rs_server</code> .	development, production
POSTGRES_TEST_DB	Test database name; defaults to <code>rs_server_test</code> .	test
SMTP_USERNAME, SMTP_PASSWORD	Credentials for the configured SMTP service.	production mail
SMTP_DOMAIN_NAME, SMTP_DOMAIN_ADDRESS	Sender domain and SMTP server address.	production mail
EMAIL_BUG_REPORT, EMAIL_ADMIN	Recipient addresses for reports and general contact messages.	development, production
RAILS_LOG_TO_STDOUT	When present, sends production Rails logs to standard output.	production
RAILS_MAX_THREADS	Maximum Active Record connection pool size; defaults to 5.	all environments
PUBLIC_BASE_URL	Public HTTP or HTTPS origin used for external links and allowed host validation; required in production.	production
ADDITIONAL_ALLOWED_HOSTS	Additional accepted request hosts, separated by commas.	production
CORS_ALLOWED_ORIGINS	Origins allowed to call the API, separated by commas. An omitted value disables requests from other origins in production.	production
FORCE_SSL	Set to <code>true</code> when the public instance is served through HTTPS.	production
RS_PDF_MAX_DOWNLOAD_BYTES	Maximum accepted publication size in bytes; defaults to 52428800 (50 MiB).	development, production
RS_PDF_OPEN_TIMEOUT	Maximum time allowed to open a connection, in seconds; defaults to 5.	development, production

Environment Variable	Purpose	Scope
RS_PDF_READ_TIMEOUT	Maximum time allowed to read a response, in seconds; defaults to 20.	development, production
RS_PDF_PROCESS_TIMEOUT	Maximum RS_PDF execution time in seconds; defaults to 60.	development, production
RS_PDF_ALLOW_PRIVATE_NETWORKS	Set to true only when publications must be fetched from a trusted private network.	development, production
RS_PDF_STORAGE_ROOT	Private directory used for prepared publications; defaults to storage/publications .	development, production
RS_PDF_DOWNLOAD_URL_TTL	Lifetime of a signed publication download URL, in seconds; defaults to 300.	development, production
RS_AUTHENTICATION_RATE_LIMIT	Authentication attempts allowed for one IP address in three minutes; defaults to 10.	development, production
RS_PASSWORD_RECOVERY_IP_RATE_LIMIT	Recovery requests allowed for one IP address in fifteen minutes; defaults to 5.	development, production
RS_PASSWORD_RECOVERY_ACCOUNT_RATE_LIMIT	Recovery requests allowed for one normalized email address in thirty minutes; defaults to 3.	development, production
RS_CONTACT_RATE_LIMIT	Contact messages allowed for one IP address in ten minutes; defaults to 5.	development, production
RS_PDF_PROCESSING_RATE_LIMIT	Combined publication processing requests allowed for one reader in one hour; defaults to 30.	development, production

Table 2: Environment variables of RS_Server.

3.5.4 Setting Variables

For local and Compose deployments, copy `.env.example` to `.env` in the RS_Server root and replace the example values using `key=value` entries. Do not quote secrets unless the selected environment loader requires it, and never commit the resulting `.env` file.

The following listing shows a valid `.env` file:

```
SECRET_PROD_KEY=replace_with_a_long_random_secret
POSTGRES_USER=postgres
POSTGRES_PASSWORD=replace_with_a_database_password
POSTGRES_DB=rs_server
```

```
SMTP_USERNAME=replace_with_an_smtp_user
SMTP_PASSWORD=replace_with_an_smtp_password
SMTP_DOMAIN_NAME=readersourcing.example
SMTP_DOMAIN_ADDRESS=smtp.example
EMAIL_BUG_REPORT=bugs@readersourcing.example
EMAIL_ADMIN=contact@readersourcing.example
PUBLIC_BASE_URL=https://readersourcing.example
CORS_ALLOWED_ORIGINS=https://client.readersourcing.example
FORCE_SSL=true
RAILS_LOG_TO_STDOUT=true
```

The value of `PUBLIC_BASE_URL` is an origin: it contains the scheme, host, and optional port, but no path, query string, fragment, or credentials. Production refuses to start when the value is missing or invalid. Its host is accepted by the Rails host authorization policy; further trusted hosts can be listed through `ADDITIONAL_ALLOWED_HOSTS`. The public origin is used for password recovery and prepared publication links rather than deriving either value from an untrusted request host.

`CORS_ALLOWED_ORIGINS` contains exact origins separated by commas. Supplying `*` deliberately permits requests from every origin and is inappropriate for an ordinary public deployment. `RS_Rate` requests browser permission for the selected server origin and therefore does not depend on a generated extension origin being listed here. Set `FORCE_SSL=true` only after HTTPS is available at the public origin.

Publication downloads reject private and reserved network addresses by default. The `RS_PDF_ALLOW_PRIVATE_NETWORKS` override is intended only for a controlled deployment in which trusted publications are deliberately served from such a network.

Prepared publications are stored below `RS_PDF_STORAGE_ROOT` and not below `public`. The API retains its original publication path fields as logical compatibility metadata, while `pdf_download_url` and `pdf_download_url_link` contain signed URLs with the configured lifetime. Logging out does not revoke an issued download URL, which remains usable only until that short lifetime ends. Existing files under `public/user` can be moved once with `bin/rails publications:migrate_private_storage`; the task stops rather than combining an existing destination with the old tree.

Rate limit counters use the configured Rails cache store. The supplied Puma configuration runs a single process and uses a store in memory. Deployments with multiple server processes or replicas must configure a shared Active Support cache store so that every instance contributes to the same request budgets.

3.5.5 Sending Mails

`RS.Server` supports standard SMTP services for account confirmation, password recovery, password change notifications, bug reports, and contact messages. Recovery messages contain the link described above, which can be used once; the new password is entered only in `RS.Server` and is not included in either recovery or confirmation mail. The SMTP

provider must expose a host, port 587 with STARTTLS, a user name, a password, and a verified sender domain. API keys specific to a provider can be placed in `SMTP_PASSWORD` when the provider uses an API key as its SMTP credential; the provider documentation defines the required values.

3.5.6 Connecting To The Database

A full connection string supplied through `DATABASE_URL` takes precedence over the individual `POSTGRES_*` variables in development and production. The test environment follows the same rule using `TEST_DATABASE_URL` and `POSTGRES_TEST_DB`.

```
ENV['DATABASE_URL'] ||
  "postgresql://#{ENV['POSTGRES_USER']} | 'postgres'}:" \
  "#{ENV['POSTGRES_PASSWORD']}@" \
  "#{ENV['POSTGRES_HOST']} | 'localhost'}/" \
  "#{ENV['POSTGRES_DB']} | 'rs_server'}"
```

3.5.7 Logging To The Standard Output

Development writes logs to standard output by default. In production, setting `RAILS_LOG_TO_STDOUT` configures a tagged logger on standard output, which is appropriate for Docker and other distributed deployments.

4 RS_PDF

`RS_PDF` [6] is the software library used by `RS_Server` to edit PDF files and add the required URL when a reader requests to save a publication for later reading. The current implementation appends a final page containing a clickable rating URL and a QR code, and stores the same URL in the PDF `BaseUr1` metadata field. In this integration, `RS_PDF` treats the complete URL as opaque data: the paper rating reference is created and later validated only by `RS_Server`.

Its command line interface allows `RS_Server` to invoke it directly. Because the components are deployed together, they do not require a separate network service or communication protocol. The executable is distributed as a single shaded JAR named `RS_PDF-v2.0.0-shaded.jar` and bundled in the `RS_Server lib` directory.

4.1 Implementation and Technology

`RS_PDF` is developed in Kotlin, an object oriented programming language that interoperates with the Java Virtual Machine. This interoperability allows developers to reuse software distributed in JAR format and call Java classes directly from Kotlin. The current build targets Java 21 and uses Kotlin 2.4.

This programming language was chosen because it incorporates many modern features and receives robust support. More importantly, the underlying tool used to edit PDF files is




Figure 6: Package diagram of RS_PDF.

PDFBox,¹¹ a Java library that provides a complete PDF manipulation toolkit. Therefore, RS_PDF serves as a wrapper for PDFBox and adds the required reference to the PDFs requested by readers. ZXing¹² generates the QR code placed beside the link annotation, Apache Commons CLI parses the execution parameters, and Log4j 2 provides logging.

Kotlin was created by JetBrains,¹³ which announced full support for Android development with Google in 2017.¹⁴ In the same year, JetBrains also announced Kotlin/Native, which compiles Kotlin programs to native binaries without requiring the JVM.

The Kotlin documentation includes an official comparison with Java,¹⁵ and the language has been discussed extensively by software developers.¹⁶

4.2 Package Diagram

Figure 6 shows how RS_PDF is divided into packages and provides a general overview of its internal architecture.

Interaction with RS_Server begins in the **program** package. The *Program* object parses and validates the command line options before delegating to the publication controller.

The **utils** package provides shared filesystem defaults, program constants, and logging configuration. Server addresses are supplied through the required URL argument and are not embedded as constants in the code.

The **publications** package contains the PDF business logic. Its Controller/Model separation follows the same object oriented intent as MVC without depending on Rails: the controller discovers one or more PDF files and creates the corresponding models, while each

¹¹<https://pdfbox.apache.org/>

¹²<https://github.com/zxing/zxing>

¹³<https://www.jetbrains.com/>

¹⁴<https://blog.jetbrains.com/kotlin/2017/05/kotlin-on-android-now-official/>

¹⁵<https://kotlinlang.org/docs/reference/comparison-to-java.html>

¹⁶<https://medium.com/@octskyward/why-kotlin-is-my-next-programming-language-c25c001e26e3>




Figure 7: General class diagram of RS_PDF v2.0.0.

model loads a document, adds the link and QR code, writes metadata, and saves a new file. A View is unnecessary because the result is the edited PDF itself.

4.3 Class Diagram

Figure 7 shows the main classes of RS_PDF. The structure follows Controller/Model responsibilities: *Program* owns process parsing and flow, *PublicationController* manages the collection, and *PublicationModel* performs PDF transformations through PDFBox and ZXing.

The *Parameters* data class transports the caption, URL, optional authentication token, and optional publication identifier through the controller boundary. This keeps the *PublicationModel* interface stable as related input data evolve.

4.4 Installation

RS_PDF is bundled with compatible RS_Server distributions, eliminating the need for a separate installation. To use it independently, download the shaded JAR from the repository releases¹⁷ or build it from source:

```
./mvnw clean verify
```

On Windows use `mvnw.cmd clean verify`. The Maven Wrapper selects the configured Maven version and the executable is generated as `target/RS_PDF-v2.0.0-shaded.jar`.

4.4.1 Requirements

- Java runtime ≥ 21 to execute the shaded JAR;
- JDK ≥ 21 to compile and verify the project from source.

¹⁷https://github.com/Miccighel/Readersourcing-2.0-RS_PDF

4.4.2 Command Line Interface

The behavior of RS_PDF is configured through the command line options in Table 3. Caption and URL are always required. Input and output paths may be omitted only together, in which case the default `data` and `res` directories are used. The authentication token and publication identifier are optional, but if either is supplied then both custom paths and both server values must be supplied.

Short	Long	Description	Constraint
<code>-pIn</code>	<code>--pathIn</code>	Input PDF file or directory, using a relative or absolute path.	Optional; requires <code>-pOut</code> .
<code>-pOut</code>	<code>--pathOut</code>	Existing output directory for annotated PDF files.	Optional; requires <code>-pIn</code> .
<code>-c</code>	<code>--caption</code>	Caption of the clickable link added to the PDF.	Required.
<code>-u</code>	<code>--url</code>	Rating URL used by the link, QR code, and <code>BaseUrl</code> metadata.	Required valid URL.
<code>-a</code>	<code>--authToken</code>	Authentication token received from RS_Server.	Optional; requires paths and <code>-pId</code> .
<code>-pId</code>	<code>--publicationId</code>	Publication identifier received from RS_Server.	Optional; requires paths and <code>-a</code> .

Table 3: Command line options of RS_PDF.

The following example assumes that:

- there is a folder containing n PDF files at `C:\data`;
- annotated files must be saved in the existing folder `C:\out`;
- the shaded JAR is located at `C:\lib\RS_PDF-v2.0.0-shaded.jar`.

RS_PDF can then be executed with the following command:

```
java -jar C:\lib\RS_PDF-v2.0.0-shaded.jar \  
-pIn C:\data -pOut C:\out \  
-c "Express your rating" \  
-u "https://readersourcing.example/rate"
```

5 RS_Rate

RS_Rate [8] is an extension designed to function as a client for the Readersourcing 2.0 ecosystem without requiring access to its website. Compatible with Google Chrome,¹⁸ Microsoft Edge,¹⁹ and Mozilla Firefox,²⁰ the extension allows users to rate publications directly from their browsers. This eliminates the need to navigate to the main website, streamlining the process of providing ratings for publications.

The primary objective of RS_Rate is to enable readers to rate a publication with minimal effort: just a few clicks or keystrokes within the Readersourcing 2.0 ecosystem. In doing so, it contributes to a more dynamic online reading experience. RS_Rate is the first client developed for the project beyond the web interface available on the main portal.

Looking ahead, our vision includes expanding the compatibility of RS_Rate to Safari and other popular browsers. Our commitment is to make this rating tool accessible across a broad range of browsers, ensuring users can seamlessly interact with content and provide feedback regardless of their preferred browser. The current Chromium and Firefox targets share the same application source and rating logic.

5.1 Implementation and Technology

RS_Rate is an extension for Google Chrome, Microsoft Edge, and Mozilla Firefox. These extensions are developed using standard web technologies such as HTML, CSS, and JavaScript. They are collections of files packaged according to the requirements of each browser. The build process copies and checks all executable assets locally, then produces unpacked Manifest V3 extensions in `dist/chromium` and `dist/firefox`.

The JavaScript component also uses jQuery to simplify DOM selection and manipulation, event management, animations, and AJAX operations. RS_Rate relies on these AJAX features to provide responsive interactions with the server. Bootstrap 4 and its plugins based on jQuery provide the markup and interaction behavior used by the extension views. Dependencies remain declared in `package.json` and are resolved through the committed Yarn lockfile.

5.2 Installation

The original 2019 release of RS_Rate remains available on the Google Chrome Web Store. The current development source produces the Manifest V3 packages described in the following section. To install the original release with a browser based on Chromium, use the store page:

- **Google Chrome and Microsoft Edge:** <https://chromewebstore.google.com/detail/readersourcing-20-rsrate/hlkdlngpijhdkbdhlmgeemffaocjagg>

¹⁸<https://www.google.com/chrome/>

¹⁹<https://www.microsoft.com/en-us/edge/>

²⁰<https://www.mozilla.org/firefox/>

The repository also produces a Firefox Manifest V3 package ready for local validation and submission to Mozilla Add-ons. It uses the following stable Gecko identifier:

```
rs-rate@readersourcing.org
```

For local testing, open `about:debugging`, choose *This Firefox*, and load

```
dist/firefox/manifest.json
```

using *Load Temporary Add-on*. Distribution through Mozilla Add-ons requires validation and signing by Mozilla.

5.3 Development and Packaging

The development toolchain uses Node.js 24 LTS and Yarn 4.17.1. From the repository root:

```
corepack enable
yarn install --immutable
yarn verify
```

`yarn verify` builds both targets, checks that all runtime assets are local, validates the Firefox build with Mozilla's `web-ext`, and runs the client contract tests. The Chromium and Firefox archives are generated with:

```
yarn package
```

The files `rs_rate-chromium.zip` and `rs_rate-firefox.zip` are written to `artifacts`. The commands `yarn package:chromium` and `yarn package:firefox` create the individual archives. The Firefox target requires Firefox 142 or later. Its manifest declares the account, browsing, website interaction, and publication content categories required by the Readersourcing workflow. The extension does not include an analytics channel.

Development builds initially target `http://localhost:3000`. Users can change the `RS_Server` host in the extension options, and the same configured host is used for normal API requests and PDF extraction. When the value is saved, the browser asks the reader to grant access to that HTTP or HTTPS origin. The permission covers only the selected origin, while any configured application path remains part of the server address.

5.4 Usage

Figure 8 shows a Google Chrome window with the extension active for a publication. This represents a typical reader visiting a publisher's website to access a paper in PDF format. The figure also shows the first page displayed when the reader opens the client. From this page, the reader can reach either the login page shown in Figure 9 or the registration page. The login page also provides access to a password recovery page, which requests the email containing the link to the `RS_Server` password form.




Figure 8: RS_Rate characterized as an extension having a popup action.

The image shows the login page of RS_Rate. It features a blue header bar with a back arrow and a settings gear. Below the header, the title "READERSOURCING 2.0" is displayed. The page prompts the user to "Insert your login data" and provides input fields for "Email (e.g. address@example.com)" and "Password". There are links for "Did you forget your password?" and "Do you want to create a new account?". A green "LOGIN" button is at the bottom.

Figure 9: The login page of RS_Rate.




Figure 10: The rating page of RS_Rate.

If a reader has yet to register for Readersourcing 2.0, they can navigate from the page shown in Figure 8 to the one displayed in Figure 11 and fill in the registration form. Once they complete the standard registration and login operations, they will find themselves on the rating page, as shown in Figure 10.

In the central section of the rating page, a reader can use the slider to choose a value in the 0–100 interval. After selecting the desired value, they submit the rating by clicking the green *Rate* button. The options page also allows the reader to make the rating anonymous. Readers must still be logged in when submitting an anonymous rating so that the system can prevent spam and repeated ratings of the same publication. Apart from these checks, the rating is processed without exposing the reader’s identity.

If the reader prefers to rate the publication later, they can use the *Download* button in the *Save For Later* section. Before the operation begins, the extension reports that the publication is ready to prepare. During the request, the same status area reports validation, annotation, and final verification. Once the procedure is complete, the prepared publication can be opened, as shown in Figure 12.

If the publication server requires access through the reader’s browser session or cannot be reached, the interface explains the condition and displays the local PDF alternative. The reader can upload the original file already present on the device without changing the publication URL used by the domain model. This source follows the same validation and verification procedure as a direct retrieval. Other conditions, including size limits, encrypted or malformed PDFs, annotation failures, and verification failures, are reported separately




Figure 11: The user registration page of RS_Rate.

instead of being presented as a generic download error.

The reader can then download the annotated publication. The refresh button, located to the right of the *Download* button, retrieves a new copy of the publication, annotates it, and makes it available again. This feature is useful when an author updates a publication after the first download.

The downloaded PDF contains a new final page with the rating URL. The URL includes the encrypted paper rating reference associated with the reader and publication, not the reader's JWT. The reference has no automatic expiration and remains independent of login and logout operations. Figure 13 shows an example opened in a PDF reader.

Following the reference takes the reader to a dedicated page in RS_Server. This interface allows them to rate the publication independently of the browser extension used to store the reference. For example, a reader can transfer the annotated publication to a tablet and use its integrated browser to submit the rating. Figure 14 shows the interface displayed after the reader follows the stored reference. The reader must have a current authenticated session, so another person who receives a copy of the publication cannot use the reference on their behalf. If the reader has logged out, they sign in again and reopen the same reference; the PDF itself does not need to be annotated again.

Whenever a reader rates a publication, RS_Server updates the scores computed by both the RSM and TRM models. RS_Rate displays the two publication scores at the bottom of the rating page, as shown in Figures 10 and 12. The reader can view their own scores by




Figure 12: The rating page of RS_Rate after a “save for later” request.




Figure 13: A publication annotated through RS_Rate.

The image shows a web interface for 'READERSOURCING 2.0'. At the top, it says 'Rate this publication'. Below this is a horizontal slider with a blue dot in the middle, and the number '50' is displayed below the slider. Under the slider, there is a checkbox labeled 'Anonymize' which is currently unchecked. Below that is a link that says 'Validate yourself'. Then there is a text input field for 'Email (e.g. address@example.com)' containing the text 'johndoe@mail.com'. Below the email field is a 'Password' label and a password input field with six dots. At the bottom of the form is a green button with the word 'RATE' in white capital letters.

Figure 14: The server interface to rate a publication.

clicking the profile button in the upper right corner. The profile page shown in Figure 15 also allows the reader to change their password.

6 RS_Py

RS_Py [7] is an additional component of the Readersourcing 2.0 ecosystem, providing a fully functional Python implementation of the models presented by Soprano et al. [5]. These models are also encapsulated by the server application of Readersourcing 2.0, as described in Section 3.

Developers familiar with Python can use RS_Py to generate and test new simulations of readers rating a set of publications. They can alter the internal logic of the models to test new approaches without having to fork and edit the full Readersourcing 2.0 implementation.

6.1 Implementation and Technology

RS_Py is a collection of *Jupyter Notebooks*. Jupyter²¹ is an open source web application powered by Python²² for creating and sharing documents containing live code, equations,

²¹<https://jupyter.org/>

²²<https://www.python.org/>




Figure 15: The profile page of RS_Rate.

visualizations, and narrative text. Notebooks are composed of cells that can be run independently, providing a detailed overview of the implemented computations.

6.2 Installation

To use the RS_Py notebooks, clone the repository²³ and place it somewhere on the filesystem. The project targets Python 3.13, and its tested direct dependencies are declared in `requirements.txt`. An isolated environment can be created as follows:

```
python -m venv .venv
source .venv/bin/activate
python -m pip install -r requirements.txt
```

On Windows, activate the environment with `.venv\Scripts\activate`.

6.3 Usage

RS_Py is organized into the following main folders:

- The `data` folder stores the datasets used to test the models presented by Soprano et al. [5] and the compressed synthetic datasets.
- The `models` folder stores the output of these models and the computed quantities.

²³https://github.com/Miccighel/Readersourcing-2.0-RS_Py

- The **notebooks** folder contains Jupyter notebooks used to generate datasets, implement the models, and analyze experiments.
- The **scripts** folder is generated on demand and contains implementations of the Jupyter notebooks as pure Python scripts.
- The **src** folder contains the conversion utility and the domain implementation held in memory.
- The **tests** folder contains the numerical and utility characterization tests.
- The **utils** folder contains auxiliary tooling, including the power law generator.

Within the **notebooks** folder, the principal files are:

- **1.Seed-Power-Law.ipynb** generates synthetic data using power law distributions;
- **2.Readersourcing.ipynb** implements and runs the RSM model;
- **3.TrueReview.ipynb** implements and runs the TRM model;
- **4.Experiments.ipynb** supports experimental analysis;
- **ReadersourcingToolkit.py** contains shared Python functionality used by the notebooks.

The **src/Convert.py** script converts the numbered notebooks into pure Python scripts and stores them in the generated **scripts** folder together with their shared toolkit.

The behavior of the notebooks can be customized by modifying the parameter settings found in their initial cells. To run them, navigate to the **notebooks** folder and type **jupyter lab**. This command starts the Jupyter server and opens the interface from which the notebooks can be executed in numerical order.

The implementation held in memory under **src/rs_py** and its characterization tests reproduce the reader, publication, and rating behavior currently shared with RS_Server. They complement rather than restrict the principal purpose of RS_Py: developers remain free to alter the internal logic of the notebooks and test new approaches. Additional object oriented simulation work, including an explicit author entity, is maintained in the Readersourcing_OO repository.²⁴

7 Overview

Figure 16 summarizes the capabilities of Readersourcing 2.0. In this example, four readers—RD1, RD2, RD3, and RD4—use RS_Rate to rate publication P1. Readers RD1, RD2, and RD3 choose *Save for later* and receive a version of the publication annotated with a link, denoted by P1+Link.

²⁴https://github.com/EddyMaddalena/Readersourcing_OO




Figure 16: Readers' interaction modalities with the Readersourcing 2.0 ecosystem.

After some time, RD1 opens P1+Link with a preferred PDF reader. RD2 sends it to a tablet, while RD3 opens it in a desktop browser. By following the link or scanning the QR code added through RS_PDF, they reach the dedicated RS_Server page and provide their ratings. In contrast, RD4 rates P1 immediately through the RS_Rate interface.

References

- [1] Luca De Alfaro and Marco Faella. TrueReview: A Platform for Post-Publication Peer Review. In: *CoRR* abs/1608.07878 (2016). arXiv: 1608.07878. URL: <http://arxiv.org/abs/1608.07878>.
- [2] Martin Fowler. *UML Distilled: A Brief Guide to the Standard Object Modeling Language*. 3rd ed. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc., 2003. ISBN: 0321193687.
- [3] Erich Gamma et al. *Design Patterns: Elementi per il Riutilizzo di Software ad Oggetti*. Pearson, Jan. 2002.
- [4] Stefano Mizzaro. Quality control in scholarly publishing: A new proposal. In: *Journal of the American Society for Information Science and Technology* 54.11 (2003), pp. 989–1005. DOI: 10.1002/asi.10296. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/asi.10296>.
- [5] Michael Soprano and Stefano Mizzaro. Crowdsourcing Peer Review: As We May Do. In: *Digital Libraries: Supporting Open Science*. Vol. 988. Communications in Computer and Information Science. Springer, 2019, pp. 259–273. DOI: 10.1007/978-3-030-11226-4_21. URL: https://doi.org/10.1007/978-3-030-11226-4_21.
- [6] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS_PDF*. DOI: 10.5281/zenodo.1442597. URL: <https://doi.org/10.5281/zenodo.1442597>.
- [7] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS_Py*. DOI: 10.5281/zenodo.3245208. URL: <https://doi.org/10.5281/zenodo.3245208>.
- [8] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS_Rate*. DOI: 10.5281/zenodo.1442599. URL: <https://doi.org/10.5281/zenodo.1442599>.
- [9] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS_Server*. DOI: 10.5281/zenodo.1442630. URL: <https://doi.org/10.5281/zenodo.1442630>.
- [10] Michael Soprano, Kevin Roitero, and Stefano Mizzaro. HITS Hits Readersourcing: Validating Peer Review Alternatives Using Network Analysis. In: *Proceedings of the 4th Joint Workshop on Bibliometric-enhanced Information Retrieval and Natural Language Processing for Digital Libraries*. Vol. 2414. CEUR Workshop Proceedings. 2019, pp. 70–82. URL: <https://ceur-ws.org/Vol-2414/paper7.pdf>.
- [11] Michael Soprano et al. Evaluation of Crowdsourced Peer Review using Synthetic Data and Simulations. In: *Proceedings of the 21st Conference on Information and Research Science Connecting to Digital and Library Science*. Vol. 3937. CEUR Workshop Proceedings. 2025. URL: <https://ceur-ws.org/Vol-3937/paper8.pdf>.