

UNIVERSITY OF UDINE

---

DEPARTMENT OF MATHEMATICS, COMPUTER SCIENCE AND PHYSICS

READERSOURCING 2.0:  
DOCUMENTATION  
VERSION 2.0.0

MICHAEL SOPRANO AND STEFANO MIZZARO

---

# Contents

|                                                |            |
|------------------------------------------------|------------|
| <b>List of Figures</b>                         | <b>iii</b> |
| <b>List of Tables</b>                          | <b>iv</b>  |
| <b>1 Introduction</b>                          | <b>1</b>   |
| <b>2 General Architecture</b>                  | <b>1</b>   |
| <b>3 RS_Server</b>                             | <b>1</b>   |
| 3.1 Implementation and Technology . . . . .    | 2          |
| 3.2 Communication Paradigm . . . . .           | 3          |
| 3.3 Database . . . . .                         | 6          |
| 3.4 Class Diagram . . . . .                    | 7          |
| 3.5 Deploy . . . . .                           | 14         |
| 3.5.1 Modality 1: Manual . . . . .             | 14         |
| 3.5.1.1 Requirements . . . . .                 | 14         |
| 3.5.1.2 How To . . . . .                       | 14         |
| 3.5.1.3 Quick Cheatsheet . . . . .             | 15         |
| 3.5.2 Modality 2: Docker Compose . . . . .     | 15         |
| 3.5.2.1 Requirements . . . . .                 | 15         |
| 3.5.2.2 How To . . . . .                       | 16         |
| 3.5.3 Environment Variables . . . . .          | 16         |
| 3.5.4 Setting Variables . . . . .              | 19         |
| 3.5.5 Sending Mails . . . . .                  | 20         |
| 3.5.6 Connecting To The Database . . . . .     | 20         |
| 3.5.7 Logging To The Standard Output . . . . . | 21         |
| <b>4 RS_PDF</b>                                | <b>21</b>  |
| 4.1 Implementation and Technology . . . . .    | 21         |
| 4.2 Package Diagram . . . . .                  | 22         |
| 4.3 Class Diagram . . . . .                    | 22         |
| 4.4 Installation . . . . .                     | 23         |
| 4.4.1 Requirements . . . . .                   | 23         |
| 4.4.2 Command Line Interface . . . . .         | 23         |
| <b>5 RS_Rate</b>                               | <b>24</b>  |
| 5.1 Implementation and Technology . . . . .    | 25         |
| 5.2 Installation . . . . .                     | 25         |
| 5.3 Development and Packaging . . . . .        | 26         |
| 5.4 Usage . . . . .                            | 26         |

|          |                                         |           |
|----------|-----------------------------------------|-----------|
| <b>6</b> | <b>RS_Py</b>                            | <b>30</b> |
| 6.1      | Implementation and Technology . . . . . | 32        |
| 6.2      | Installation . . . . .                  | 32        |
| 6.3      | Usage . . . . .                         | 32        |
| <b>7</b> | <b>Overview</b>                         | <b>33</b> |
|          | <b>References</b>                       | <b>34</b> |

## List of Figures

|    |                                                                                                                                 |    |
|----|---------------------------------------------------------------------------------------------------------------------------------|----|
| 1  | Conceptual architecture of Readersourcing 2.0 (NOT UML). RS_PDF is bundled with and invoked by RS_Server. . . . .               | 2  |
| 2  | Intuitive scheme of the MVC pattern (NOT UML). . . . .                                                                          | 3  |
| 3  | Persistent entities and associations in the RS_Server database (selected attributes). . . . .                                   | 7  |
| 4  | General class diagram of RS_Server. Method lists are intentionally selective. .                                                 | 8  |
| 5  | Authentication process of the web interface. REST clients use the returned JWT through the <b>Authorization</b> header. . . . . | 13 |
| 6  | Package diagram of RS_PDF. . . . .                                                                                              | 22 |
| 7  | General class diagram of RS_PDF v2.0.0. . . . .                                                                                 | 23 |
| 8  | RS_Rate characterized as an extension having a popup action. . . . .                                                            | 27 |
| 9  | The login page of RS_Rate. . . . .                                                                                              | 27 |
| 10 | The rating page of RS_Rate. . . . .                                                                                             | 28 |
| 11 | The user registration page of RS_Rate. . . . .                                                                                  | 28 |
| 12 | The rating page of RS_Rate after a “save for later” request. . . . .                                                            | 29 |
| 13 | A publication annotated through RS_Rate. . . . .                                                                                | 30 |
| 14 | The server interface to rate a publication. . . . .                                                                             | 31 |
| 15 | The profile page of RS_Rate. . . . .                                                                                            | 31 |
| 16 | Readers’ interaction modalities with the Readersourcing 2.0 ecosystem. . . .                                                    | 34 |

## List of Tables

|   |                                                       |    |
|---|-------------------------------------------------------|----|
| 1 | Subset of the RESTful interface of RS_Server. . . . . | 5  |
| 2 | Environment variables of RS_Server. . . . .           | 19 |
| 3 | Command line options of RS_PDF. . . . .               | 24 |

# 1 Introduction

The Readersourcing 2.0 ecosystem was developed as part of a research project funded jointly by SISSA Medialab<sup>1</sup> and the University of Udine.<sup>2</sup> It was presented by Soprano et al. [5] at the IRCDL 2019 Conference.<sup>3</sup> The paper is also available on Zenodo.<sup>4</sup> The code and related documentation are available on GitHub.<sup>5</sup>

Initially, a recap of its general architecture is presented, followed by a brief description of the role and purpose of each of its components. Specific aspects, such as the technology used, the internal architecture, and the structure of the database, are also discussed. This is achieved through different types of diagrams adhering to the UML standard unless otherwise specified, drawn according to the style rules proposed by Fowler [2]. The current software requirements and operational procedures are recorded alongside the original architectural description.

## 2 General Architecture

Readersourcing 2.0 is an ecosystem composed of several components. It includes RS\_Server [9] (presented in Section 3), which gathers the ratings given by readers, and RS\_Rate [8] (presented in Section 5), which acts as a client and enables readers to rate publications. Although all operations can be carried out directly through the web interface provided by RS\_Server, RS\_PDF [6] (presented in Section 4) annotates PDF files with a rating URL and QR code on behalf of the server application. Finally, RS\_Py (presented in Section 6) provides a fully functional Python implementation of the RSM and TRM models described by Soprano et al. [5], together with the notebooks used to generate datasets and analyze experiments.

Figure 1 provides an overview of the architecture of Readersourcing 2.0, and in the following, we briefly describe these four components. We summarize the capabilities of the ecosystem in Section 7.

## 3 RS\_Server

RS\_Server [9] is the server component responsible for collecting and aggregating the ratings given by readers and applying the RSM and TRM models described by Soprano et al. [5] to compute quality scores for readers and publications. An instance of RS\_Server is deployed with a compatible RS\_PDF executable, which is shipped in the `lib` directory and invoked as a local process when a publication must be annotated. Up to  $n$  browsers and

---

<sup>1</sup><https://medialab.sissa.it/>

<sup>2</sup><https://www.uniud.it/>

<sup>3</sup><https://ircdl2019.isti.cnr.it/>

<sup>4</sup><https://zenodo.org/record/1446468>

<sup>5</sup><https://github.com/Miccighel/Readersourcing-2.0>




Figure 1: Conceptual architecture of Readersourcing 2.0 (NOT UML). RS\_PDF is bundled with and invoked by RS\_Server.

their users can communicate with the server, each through an instance of RS\_Rate. The extension is distributed for browsers based on Chromium and for Firefox.

Readers can interact with the server either through clients installed in their browsers or through the standalone web interface provided by RS\_Server. These clients manage reader registration and authentication, rating actions, and downloads of annotated publications.

During the design phase of RS\_Server, strategies were adopted to ensure extensibility and generality. This includes:

- (i) Straightforward addition of new models.
- (ii) A shared input data format among all models.
- (iii) A standard procedure for models needing to save values locally in the RS\_Server database.

### 3.1 Implementation and Technology

RS\_Server is developed using *Ruby on Rails*,<sup>6</sup> an open source web application framework built on the Ruby programming language and strongly based on the Model View Controller (MVC) architectural pattern.

The MVC pattern separates a program's control logic from its data presentation and business logic, helping developers establish a clear architecture from the initial design phase.

<sup>6</sup><https://rubyonrails.org/>




Figure 2: Intuitive scheme of the MVC pattern (NOT UML).

Figure 2 illustrates the resulting structure. It comprises three distinct entities: the *Controller*, which manages control logic; the *Model*, which encapsulates business logic; and the *View*, which presents data.

The Controller has direct access to both the Model and the View. It typically receives user input from the View, updates the internal state of the Model through its methods, and then passes the updated Model back to the View for presentation. A software system may contain several Controllers, each of which can manage multiple Model instances. In web MVC frameworks such as Rails, Controllers are commonly organized around domain resources, and a Model can be presented through more than one View.

MVC is not the only principle underlying Rails. Another central principle is “Convention Over Configuration”: the framework reduces the number of decisions required during application development by providing standard conventions that can still be changed when greater flexibility is needed. The “Rails Doctrine” describes these principles in greater detail.<sup>7</sup>

Rails is an actively developed framework used by organizations such as *GitHub* and *SoundCloud*. It is supported by an established community and extensive learning resources.

## 3.2 Communication Paradigm

An MVC framework such as Rails can support several kinds of web applications, including a *Web Service*. A Web Service exposes operations remotely by exchanging messages encoded in a standard format such as *JSON* over a transport protocol such as *HTTP*. Its interface must define the available resources and operations, together with the messages required to access them.

<sup>7</sup><https://rubyonrails.org/doctrine/>

One communication paradigm for Web Services is *REST* (*REpresentational State Transfer*). In a RESTful interface, resources are identified by URIs and the HTTP method determines the operation to perform. The server returns a response in the agreed interchange format, and the client interprets that response.

RS\_Server is a Rails application exposing RESTful interfaces alongside web workflows rendered by the server. API clients exchange JSON messages over HTTP, while the same domain and authorization rules are reused by the web interface.

The REST interface is defined by the application routes, JSON views, and integration tests. A Postman collection provides request and response examples for the exposed resources:

<https://documenter.getpostman.com/view/4632696/2sBYAswWvk>.

For example, Table 1 presents a subset of the RESTful interface of RS\_Server. These operations manage publications, one of the entities in the application domain. If a user requests the “show” operation for publication 1 through the corresponding endpoint, RS\_Server returns a JSON response similar to the following example.

```
1 {
2   "id": 1,
3   "doi": "10.1140/epjc/s10052-018-6047-y",
4   "title": "Uncertainties in WIMP dark matter scattering revisited",
5   "subject": null,
6   "author": "John Ellis",
7   "creator": "Springer",
8   "producer": null,
9   "pdf_url": "https://example.org/publication.pdf",
10  "steadiness": 1.0,
11  "score_rsm": 0.71,
12  "score_trm": 0.68,
13  "created_at": "2026-01-01T12:00:00.000Z",
14  "updated_at": "2026-01-01T12:00:00.000Z",
15  "url": "http://localhost:3000/publications/1.json"
16 }
```

Publication records are shared among readers. The REST interface therefore exposes their creation and retrieval, together with dedicated domain operations such as fetching and refreshing, but does not expose generic update or deletion operations. This prevents one authenticated reader from arbitrarily rewriting or removing a shared publication and its associated ratings.

Operations that regenerate annotated files or change a reader’s subscription preference use POST. An unsubscribe link opens a confirmation page, and the preference changes only after the authenticated reader submits that page.

| Endpoint                                | HTTP Message | Operation          | Description                                                                                        |
|-----------------------------------------|--------------|--------------------|----------------------------------------------------------------------------------------------------|
| /publications.json                      | GET          | Index              | Fetches the entire collection of Publications.                                                     |
| /publications/list                      | GET          | List               | Fetches the entire collection of Publications (view rendered by the server).                       |
| /publications/1.json                    | GET          | Show               | Returns the Publication with identifier equal to 1.                                                |
| /publications/lookup.json               | POST         | Lookup             | Searches for a Publication by its URL without fetching it.                                         |
| /publications/random.json               | GET          | Random             | Returns a random Publication.                                                                      |
| /publications/1/is_rated.json           | GET          | Is Rated           | Checks if the Publication with identifier equal to 1 has been rated by at least one reader.        |
| /publications/1/is_saved_for_later.json | GET          | Is Saved For Later | Checks if the Publication with identifier equal to 1 has been saved for later by the current user. |
| /publications.json                      | POST         | Create             | Creates a new Publication.                                                                         |
| /publications/is_fetchable.json         | POST         | Is Fetchable       | Retrieves and opens the bounded source as a PDF for API inspection.                                |
| /publications/extract.json              | POST         | Extract            | Extracts the rating URL from an annotated Publication.                                             |
| /publications/fetch.json                | POST         | Fetch              | Retrieves, annotates, verifies, and publishes a Publication.                                       |
| /publications/fetch_upload.json         | POST         | Fetch Upload       | Performs the same preparation from an uploaded original PDF.                                       |
| /publications/1/refresh.json            | POST         | Refresh            | Fetches the Publication with identifier equal to 1 again.                                          |
| ...                                     | ...          | ...                | ...                                                                                                |

Table 1: Subset of the RESTful interface of RS\_Server.

### 3.3 Database

To support features such as storing annotated publications and authenticating users, RS\_Server uses the database schema illustrated in Figure 3. The Readersourcing 2.0 application domain contains three main entities:

- **Users:** represents the users of the system and stores their personal data. A user may have an optional *ORCID* and a boolean indicating whether they wish to receive emails. Additional attributes store the tokens required for operations such as password reset.
- **Ratings:** represents the ratings provided by readers for publications. Each rating has an integer score from 0 to 100, including both endpoints, and an original score that retains the first submitted value when the rating is edited. Two booleans indicate whether the rating is anonymous and whether it was edited later. A reader can provide at most one rating for a given publication; this invariant is enforced both by the application model and by a unique database constraint.
- **Publications:** represents publications known to the server, including those saved for later but not yet rated. A publication may have a *DOI*, descriptive metadata, and attributes used to manage the original and annotated PDF files.

These entities also contain attributes such as *steadiness* and *informativeness*, which store scores or parameters computed by the Readersourcing models.

Both rating scores are required and must remain within the rating scale. The application validates the supplied scores, and the database enforces their presence and range. This applies to ratings submitted through the API or an annotated PDF and to edits of existing ratings. An invalid value leaves the stored domain scores unchanged and does not send a confirmation email; JSON requests return HTTP status 422 with validation messages for the affected fields. Before installing the database constraints, the migration checks existing ratings and stops if their scores are missing or out of range, without correcting those records automatically.

An auxiliary **AuthenticationToken** entity records each active login. It belongs to one user and stores the unique JWT identifier and its expiration time, but never the encoded JWT itself. Authentication tokens support the security boundary rather than the Readersourcing score domain.

Figure 3 depicts the two mandatory associations of a rating. A user may give zero or more ratings, but each rating belongs to exactly one user. An anonymous rating still retains this association so that authentication and the constraint of one rating per user can be enforced; anonymity controls how the reader identity participates in presentation and processing, not referential integrity.

Similarly, a publication may receive zero or more ratings, while each rating belongs to exactly one publication. The zero multiplicity is required because a publication may be fetched and saved for later before any rating is submitted.




Figure 3: Persistent entities and associations in the RS\_Server database (selected attributes).

The auxiliary association permits a user to have zero or more active authentication tokens, while every token belongs to exactly one user. Removing a user also removes all their active tokens.

### 3.4 Class Diagram

Figure 4 shows the main responsibilities of the RS\_Server implementation. Rails controllers map HTTP operations to the domain models, while JSON and views rendered by the server present their results. The three persistent domain models encapsulate the Readersourcing logic and associations described in the previous section. The auxiliary *AuthenticationToken* model maintains the registry of active logins, while objects that are not persisted represent durable paper references and request rate policies. Authentication, reference validation, request throttling, and Readersourcing computations are delegated to focused objects rather than being absorbed into the controllers.

The *ApplicationController* contains shared authorization hooks for both API and server requests, while the *AuthenticationController* handles login and logout. The *Authenticator*, *Authorizer*, and *AuthenticationTokenRevoker* command objects isolate credential verification, token issuance, token validation, and revocation from HTTP concerns.

The *PaperRatingReference* value object similarly isolates the creation and validation of references stored in annotated publications. It uses authenticated encryption with a dedicated purpose and binds each reference to one reader and one publication without persisting a login credential in the PDF.




Figure 4: General class diagram of RS\_Server. Method lists are intentionally selective.

The *PublicationDownloadReference* value object authorizes access to a prepared copy without exposing its private storage path. A signed reference is bound to one reader, publication, copy generation, PDF variant, and file name, and expires after five minutes by default. This separate capability allows both browser interfaces to open a PDF in a new tab, where an **Authorization** header cannot be added, without making the storage directory public.

The *PublicationCopy* value object describes the files retained for one reader and publication. Its filenames and paths belong to that copy rather than to the shared publication record. The *PublicationCopyStore* command object stages the original and annotated PDFs in a separate generation and makes that pair current after the database transaction commits. A rollback removes the staged generation. Short file locks coordinate publication of the pair and retention of issued download references across application processes using the same private storage.

The *PdfOperationDeadline* value object describes the time remaining for a PDF operation through a monotonic clock. Retrieval, inspection, annotation, and verification share this deadline; a phase may shorten the remaining time but cannot renew it. The *PdfProcess* command object supervises the inspection and annotation processes and stops and reaps them when their time expires. The *PdfInspector* reads metadata and page count in a separate Ruby process using the locked **pdf-reader** dependency. Its result preserves metadata bytes and encodings through a typed JSON representation, without deserializing arbitrary Ruby objects.

The *RequestRateLimit* value object describes each request budget and provides opaque identities for the corresponding cache counters. Controllers apply these policies through the Rails rate limiting boundary, leaving credential verification, mail delivery, publication handling, and the Readersourcing domain objects unaware of throttling concerns.

The *RegistrationConfirmation* command object sends the email needed to activate a reader's account and reports delivery failure without changing the account or its confirmation token. Registration and the confirmation retry performed during authentication use the same object. Link generation and delivery failures are recorded without including email contents, credentials, or confirmation references in the log.

Registration attempts are limited to five for one IP address in fifteen minutes and three for one normalized email address in thirty minutes by default. Rejected attempts return HTTP status 429 with **Retry-After**; displaying the registration form and following a confirmation link do not consume these budgets. Creating an account returns HTTP status 201. If the confirmation email cannot be sent, the account remains unconfirmed and the response includes **status: "confirmation\_pending"**. Both interfaces display the accompanying message on the login page, asking the reader to sign in with the same credentials later. Correct credentials for an unconfirmed account request another confirmation email without replacing an existing token. The response is HTTP status 401 after sending the email, or 503 with **status: "confirmation\_delivery\_failed"** if delivery fails. Incorrect credentials do not send email or disclose that the account is awaiting confirmation. These attempts remain

subject to the authentication request limit, and no authentication token is issued before confirmation. Confirmation links use `PUBLIC_BASE_URL`; outside production, an omitted value uses the request origin.

For REST clients, relying only on session data held by the server to identify the authenticated user is unsuitable because the RESTful interface is *stateless*. Clients therefore attach an authentication *token* to each protected request. The workflow rendered by the server can still keep the same token in its Rails session.

Authentication establishes the identity of the reader making the request, while authorization at the resource level determines which records that reader may operate on. Profile updates, account deletion, and subscription changes are restricted to the current reader. A reader may inspect and edit only their own rating through the individual rating endpoints. Reader collections and individual reader responses expose the public academic profile—name, ORCID, and Readersourcing scores—without email or subscription preferences; the current reader obtains those private profile fields through the dedicated information endpoint.

When a user logs in, the supplied credentials are verified against the stored password digest and the account confirmation state. The credentials themselves are not placed in the token. If verification succeeds, `RS.Server` creates an active token record and issues a JWT signed with HS256 whose payload contains the user identifier, the client IP address, its unique `jti` identifier, and a standard expiration claim set to seven days after issuance.

The authentication response returns the JWT for REST clients. These clients attach it to protected requests through the `Authorization` header, which accepts both the direct token format used by `RS.Rate` and the conventional `Bearer <token>` format. The web interface instead keeps the token only in the encrypted Rails session, which is implemented through a cookie marked `HttpOnly`; browser scripts neither read the JWT nor add it to request headers. Upon receiving a protected request, the *Authorizer* verifies the JWT signature, expiration time, IP address, referenced user, and corresponding active token record before allowing the request to proceed. Figure 5 provides an intuitive illustration of this authentication process.

The Rails session identifier is renewed after successful authentication, before the token is stored, and is renewed again when the reader logs out. Logout removes the active token record for the current browser or API session and leaves the reader’s other devices authenticated. A successful password change or reset removes every active token belonging to that reader; account deletion removes the same records through the user association. The signed value remains immutable, but without its active record it can no longer authorize a protected request.

Logout changes application state and is therefore available only through an HTTP `POST` request. HTML requests that change state require a valid Rails authenticity token. JSON requests from the web interface use the Rails session and require the same authenticity token. JSON requests that do not rely on that session remain stateless; protected API requests carry the JWT through `Authorization` and are handled through the token authentication boundary.

In the save for later workflow, RS\_Server places a dedicated paper rating reference in the URL supplied to RS\_PDF. Its encrypted payload contains the reader and publication identifiers and has no automatic expiration. It is separate from the JWT, which is valid for seven days, and cannot authorize an API request. Following the reference always requires a current authenticated session, after which the server verifies that the reference belongs to both the selected publication and the authenticated reader. Logout, password changes, and password resets therefore do not invalidate an annotated PDF. Removing its reader or publication makes the reference unusable.

The web interface and RS\_Rate do not download a publication merely because a reader enters a URL or opens the extension. They first report that the source is ready to prepare, and the real operation starts when the reader requests the prepared PDF. RS\_Server then retrieves the source once under the configured address, redirect, timeout, and size policies. The declared content type is retained as an indication, while the file header and the PDF parser determine whether the received content is a usable PDF. Private and reserved network addresses remain refused unless a controlled deployment enables the documented override.

The same preparation object is used for remote and uploaded sources. An upload provides an alternative when a publisher requires access through the reader's browser session or when the remote source is otherwise unavailable. The source is opened before RS\_PDF is invoked and a PDF that is malformed, encrypted, unsupported, or already contains a **BaseUrl** rating reference is refused with a specific status. RS\_PDF writes into temporary storage. Before the generated file replaces any existing reader copy, RS\_Server opens it again and verifies that it contains exactly one additional page and the expected rating URL. A failed process or final verification leaves the previous copy untouched. Completed copies remain outside the public asset directory and are served only after a download reference has been verified. The response forbids shared caching and prevents the reference from being sent as a referrer. Both user interfaces translate the stable preparation statuses into an explicit title, explanation, retry action, and upload option where applicable.

Each successful preparation retains the exact original bytes together with the verified annotated PDF. The two files have names and paths specific to the reader's copy. Downloads identify a particular generation, so an issued URL continues to select the same bytes after a refresh even when the filename is reused. Subsequent preparations remove older retired generations after their issued references and a short retention margin have expired; account removal deletes every copy belonging to the reader. Original and annotated download fields are null when the corresponding file is unavailable. A publication without an annotated copy for the current reader reports **preparation\_status: "not\_prepared"**. Existing annotated files remain accessible through their signed references; if no original was retained, preparing the source again provides that file. A previous filename can be recovered from the reader's own directory when the annotated copy can be selected unambiguously.

PDF operations have a shared time budget of 120 seconds by default. Retrieval is also limited to 30 seconds, including DNS resolution, redirects, and continuous response streaming; each inspection process is limited to 15 seconds, and RS\_PDF execution to 60 seconds.

Each phase uses the shorter of its own limit and the remaining shared time. A timeout returns HTTP status 422 with `status: "processing_timeout"`, removes temporary files, and leaves any previous prepared pair available. Both interfaces explain the stopped operation and offer a retry or an original PDF upload. Process termination allows at most two additional seconds before a forced stop; on Unix, descendants in the same process group are also stopped. Process output is bounded to 1 MiB per stream. Uploaded files use this budget after Rails has received the multipart request; incoming transfer size and duration must therefore also be bounded by the HTTP proxy or hosting ingress.

To recover a rating link from an annotated publication, the reader uploads the PDF to the extraction endpoint. The upload follows the configured size limit and is opened as a PDF before its `BaseUrl` metadata is read. The address must use the exact public origin and rating route of `RS_Server`, without credentials, a query, or a fragment. Its encrypted reference must identify an existing publication and the authenticated reader. Invalid files or references return HTTP status 422, and the temporary copy is removed when the request finishes. The web interface also verifies the origin and route before opening the returned link and displays an error if validation fails.

Password recovery uses a separate credential that can be used once. An HTTP `GET` request only displays the recovery form; sending the request requires `POST`. The request endpoint returns the same confirmation for registered and unregistered addresses. For a registered address, `RS_Server` generates a random token, stores only its SHA-256 digest and issue time, and sends a link built from the configured public origin. The link remains valid for four hours. It opens a form rendered by the server in which the reader chooses and confirms a new password; a successful reset clears the stored digest before a confirmation email is sent. Passwords are never sent by email.

Request rate policies bound operations that can be abused or consume disproportionate resources. By default, one IP address may make ten authentication attempts in three minutes and five password recovery requests in fifteen minutes; one normalized email address may receive three recovery requests in thirty minutes; one IP address may send five contact messages in ten minutes; and one authenticated reader may initiate thirty combined publication download, inspection, extraction, and annotation operations in one hour. Exceeding a budget returns HTTP status 429 `Too Many Requests` with a `Retry-After` header and does not execute the protected action. Cache identities are SHA-256 digests rather than plain email addresses or IP values. Displaying the password recovery form does not consume a recovery request.

Every HTML and JSON response carries an enforced Content Security Policy. Browser scripts, styles, and fonts are installed through Yarn and served by the Rails asset pipeline, so executable code and presentation resources do not depend on external hosts at request time. Direct versions are declared in `package.json` and the complete dependency graph is recorded in `yarn.lock`. Visual dependencies remain within the compatibility lines used by the original interface. `DataTables`, `JSZip`, and `js-cookie` use the first versions that address their known registry advisories while preserving the APIs used by `RS_Server`. `Pdfmake` uses

the 0.2 compatibility line, which preserves the API required by the original DataTables integration. Pdfmake and its font data are loaded only on the authenticated publication and reader lists, where they provide the original PDF export behaviour without allowing dynamic code evaluation. The original views and visual assets require style attributes, which remain permitted. Images are restricted to the application, data URLs, and the badge host used on the resources page. Other directives forbid object content, external form targets, external base URLs, and framing.

Responses set `X-Frame-Options` to `DENY`, `X-Content-Type-Options` to `nosniff`, and `Referrer-Policy` to `same-origin`. The `Permissions-Policy` header disables access to the camera, screen capture, location, microphone, payment, and USB interfaces. Rails adds HTTP Strict Transport Security only when `FORCE_SSL=true`.




Figure 5: Authentication process of the web interface. REST clients use the returned JWT through the `Authorization` header.

The classes responsible for the Readersourcing computations follow the *Strategy* pattern: *Readersourcing*, *ReadersourcingStrategy*, *RsmStrategy*, and *TrmStrategy*. A rating invokes the RSM strategy first and the TRM strategy second. The pattern keeps model selection extensible while expressing the score computation order explicitly.

The *RatingSubmission* command object coordinates the persistence of ratings and the results produced by these strategies. It saves a new rating and the RSM and TRM results in one database transaction. Rating submissions and edits acquire the same PostgreSQL advisory transaction lock: reader reputation is shared across publications, so locking only the selected publication would not protect every affected score. The reader and publication are read again after the lock is acquired. Each strategy requires its writes to succeed; a

failed calculation rolls back the rating, any publication created for it, and all associated score changes. The lock coordinates rating operations across application processes and is released when their transaction ends. Editing a rating retains its original score and does not recompute either model.

Confirmation mail is attempted for subscribed readers only after the new rating and its calculations have been committed. A mail delivery failure leaves the saved rating intact and does not change the successful API response. The paper rating interface reports that the rating was saved but its confirmation email could not be sent.

## 3.5 Deploy

RS.Server supports two setup paths: a manual installation, intended primarily for development, and a Docker Compose deployment, which provides a reproducible application and database environment. Docker Compose can use the published `miccighel/rs_server:v2.0.0` image available on Docker Hub,<sup>8</sup> or build the application locally from the same source tree.

Secrets and database settings specific to each environment are described in Section 3.5.3. A local `.env` file must never be committed.

### 3.5.1 Modality 1: Manual

This modality downloads and initializes RS.Server directly on the local machine and offers the greatest flexibility for development and domain model experiments.

#### 3.5.1.1 Requirements

- Ruby == 3.4.10;
- Bundler compatible with the committed `Gemfile.lock`;
- a Java 21 runtime, required by RS.PDF;
- PostgreSQL >= 17;
- Node.js >= 18, used to install and precompile the existing browser assets.

#### 3.5.1.2 How To

Clone the RS.Server repository,<sup>9</sup> navigate to its root directory, copy `.env.example` to `.env` as described in Section 3.5.4, and run:

```
bundle install
node .yarn/releases/yarn-3.6.3.cjs install --immutable
bin/rails db:create
```

---

<sup>8</sup>[https://hub.docker.com/r/miccighel/rs\\_server](https://hub.docker.com/r/miccighel/rs_server)

<sup>9</sup>[https://github.com/Miccighel/Readersourcing-2.0-RS\\_Server](https://github.com/Miccighel/Readersourcing-2.0-RS_Server)

```
bin/rails db:migrate
bin/rails db:seed # optional sample data
bin/rails server -b 127.0.0.1 -p 3000 -e development
```

PostgreSQL must be accepting connections before the database commands are executed. With the sample bind address and port, the application is available at <http://127.0.0.1:3000>. Rails binstubs are always invoked from the repository root so that the locked application dependencies are used.

The active token registry is populated at login. A session carrying a JWT without a registered identifier is rejected and the reader signs in again.

### 3.5.1.3 Quick Cheatsheet

1. change to the main directory with `cd`;
2. copy `.env.example` to `.env` and populate it;
3. `bundle install`;
4. `node .yarn/releases/yarn-3.6.3.cjs install --immutable`;
5. `bin/rails db:create`;
6. `bin/rails db:migrate`;
7. `bin/rails db:seed` (optional);
8. `bin/rails server -b <address> -p <port> -e development`.

## 3.5.2 Modality 2: Docker Compose

Docker Compose uses the public `miccighel/rs_server:v2.0.0` image and starts it together with PostgreSQL 17. The application image includes Ruby 3.4.10, the Java 21 runtime required by RS\_PDF, the locked Ruby dependencies, and the precompiled browser assets. When the source code has been edited, the same image can be built locally. This is the recommended reproducible deployment for local integration and verification under conditions similar to production. Named volumes retain the database and the prepared publications when the containers are recreated. The application runs as an unprivileged user with a read only root file system; temporary directories and the private publication volume are its only writable locations.

### 3.5.2.1 Requirements

- Docker Desktop or another installation supporting Docker Compose v2.<sup>10</sup>

---

<sup>10</sup><https://docs.docker.com/compose/>

### 3.5.2.2 How To

Clone the repository, copy `.env.example` to `.env` and replace its example values, ensure that the Docker Engine is running, and execute:

```
docker compose up
```

The command downloads the published image when it is not already available. To build the image from the checked out source instead, execute `docker compose up --build`.

The `database` service must pass its health check before the `rs_server_setup` service runs `bin/rails db:migrate`. The setup service must complete successfully before `rs_server_webapp` starts; database migration is not performed by the application entry-point and sample data are not created implicitly. Once startup completes, the web application is available at `http://localhost:3000`; `0.0.0.0` is the container bind address, not the address users should enter in a browser. The `GET /up` endpoint can be used to verify that Rails boots successfully.

Sample data remain optional:

```
docker compose run --rm rs_server_webapp bin/rails db:seed
```

Use `docker compose down` to stop the services. The named volumes are retained so that database records and prepared publications survive ordinary container recreation.

### Quick Cheatsheet

- change to the main directory with `cd`;
- copy `.env.example` to `.env` and populate it;
- `docker compose up`;
- `docker compose up --build` when the source code has been edited;
- `docker compose run --rm rs_server_webapp bin/rails db:seed` (optional);
- `docker compose down` (to stop the services).

### 3.5.3 Environment Variables

Credentials and connection details specific to each deployment are supplied through environment variables. Table 2 lists the variables used by the application. Database URLs take precedence over individual PostgreSQL settings.

| Environment Variable | Purpose                           | Scope       |
|----------------------|-----------------------------------|-------------|
| SECRET_DEV_KEY       | Rails secret used in development. | development |

| Environment Variable                     | Purpose                                                                                                                                                                      | Scope                     |
|------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------|
| SECRET_PROD_KEY                          | Rails secret used to sign and encrypt production data, including JWTs, session data, and paper rating references. Keep it stable while issued references must remain usable. | production                |
| DATABASE_URL                             | Complete PostgreSQL connection URL. Overrides the individual <b>POSTGRES_*</b> variables.                                                                                    | development, production   |
| TEST_DATABASE_URL                        | Complete PostgreSQL connection URL for the test database.                                                                                                                    | test                      |
| POSTGRES_USER                            | PostgreSQL user name used when a full URL is not supplied.                                                                                                                   | all database environments |
| POSTGRES_PASSWORD                        | Password for the PostgreSQL user.                                                                                                                                            | all database environments |
| POSTGRES_HOST                            | PostgreSQL host; defaults to <b>localhost</b> outside Compose and is set to <b>database</b> by Compose.                                                                      | all database environments |
| POSTGRES_DB                              | Development or production database name; defaults to <b>rs_server</b> .                                                                                                      | development, production   |
| POSTGRES_TEST_DB                         | Test database name; defaults to <b>rs_server_test</b> .                                                                                                                      | test                      |
| SMTP_USERNAME,<br>SMTP_PASSWORD          | Credentials for the configured SMTP service.                                                                                                                                 | production mail           |
| SMTP_DOMAIN_NAME,<br>SMTP_DOMAIN_ADDRESS | Sender domain and SMTP server address.                                                                                                                                       | production mail           |
| EMAIL_BUG_REPORT,<br>EMAIL_ADMIN         | Recipient addresses for reports and general contact messages.                                                                                                                | development, production   |
| RAILS_LOG_TO_STDOUT                      | When present, sends production Rails logs to standard output.                                                                                                                | production                |
| RAILS_MAX_THREADS                        | Maximum Active Record connection pool size; defaults to 5.                                                                                                                   | all environments          |
| PUBLIC_BASE_URL                          | Public HTTP or HTTPS origin used for external links and allowed host validation; required in production.                                                                     | production                |
| ADDITIONAL_ALLOWED_HOSTS                 | Additional accepted request hosts, separated by commas.                                                                                                                      | production                |

| Environment Variable          | Purpose                                                                                                                    | Scope                   |
|-------------------------------|----------------------------------------------------------------------------------------------------------------------------|-------------------------|
| CORS_ALLOWED_ORIGINS          | Origins allowed to call the API, separated by commas. An omitted value disables requests from other origins in production. | production              |
| FORCE_SSL                     | Set to <b>true</b> when the public instance is served through HTTPS.                                                       | production              |
| RS_PDF_MAX_DOWNLOAD_BYTES     | Maximum accepted publication size in bytes; defaults to 52428800 (50 MiB).                                                 | development, production |
| RS_PDF_OPEN_TIMEOUT           | Maximum time allowed to open a connection, in seconds; defaults to 5.                                                      | development, production |
| RS_PDF_READ_TIMEOUT           | Maximum time allowed to read a response, in seconds; defaults to 20.                                                       | development, production |
| RS_PDF_DOWNLOAD_TIMEOUT       | Maximum total retrieval time, including DNS, redirects, and streaming; defaults to 30 seconds.                             | development, production |
| RS_PDF_INSPECTION_TIMEOUT     | Maximum time for each PDF inspection process; defaults to 15 seconds.                                                      | development, production |
| RS_PDF_TOTAL_TIMEOUT          | Shared PDF operation time budget; defaults to 120 seconds.                                                                 | development, production |
| RS_PDF_PROCESS_TIMEOUT        | Maximum RS_PDF execution time in seconds; defaults to 60.                                                                  | development, production |
| RS_PDF_ALLOW_PRIVATE_NETWORKS | Set to <b>true</b> only when publications must be fetched from a trusted private network.                                  | development, production |
| RS_PDF_STORAGE_ROOT           | Private directory used for prepared publications; defaults to <b>storage/publications</b> .                                | development, production |
| RS_PDF_DOWNLOAD_URL_TTL       | Lifetime of a signed publication download URL, in seconds; defaults to 300.                                                | development, production |
| RS_AUTHENTICATION_RATE_LIMIT  | Authentication attempts allowed for one IP address in three minutes; defaults to 10.                                       | development, production |
| RS_REGISTRATION_IP_RATE_LIMIT | Registration attempts allowed for one IP address in fifteen minutes; defaults to 5.                                        | development, production |

| Environment Variable                    | Purpose                                                                                          | Scope                   |
|-----------------------------------------|--------------------------------------------------------------------------------------------------|-------------------------|
| RS_REGISTRATION_ACCOUNT_RATE_LIMIT      | Registration attempts allowed for one normalized email address in thirty minutes; defaults to 3. | development, production |
| RS_PASSWORD_RECOVERY_IP_RATE_LIMIT      | Recovery requests allowed for one IP address in fifteen minutes; defaults to 5.                  | development, production |
| RS_PASSWORD_RECOVERY_ACCOUNT_RATE_LIMIT | Recovery requests allowed for one normalized email address in thirty minutes; defaults to 3.     | development, production |
| RS_CONTACT_RATE_LIMIT                   | Contact messages allowed for one IP address in ten minutes; defaults to 5.                       | development, production |
| RS_PDF_PROCESSING_RATE_LIMIT            | Combined publication processing requests allowed for one reader in one hour; defaults to 30.     | development, production |

Table 2: Environment variables of RS\_Server.

### 3.5.4 Setting Variables

For local and Compose deployments, copy `.env.example` to `.env` in the RS\_Server root and replace the example values using `key=value` entries. Do not quote secrets unless the selected environment loader requires it, and never commit the resulting `.env` file.

The following listing shows a valid `.env` file:

```
SECRET_PROD_KEY=replace_with_a_long_random_secret
POSTGRES_USER=postgres
POSTGRES_PASSWORD=replace_with_a_database_password
POSTGRES_DB=rs_server
SMTP_USERNAME=replace_with_an_smtp_user
SMTP_PASSWORD=replace_with_an_smtp_password
SMTP_DOMAIN_NAME=readersourcing.example
SMTP_DOMAIN_ADDRESS=smtp.example
EMAIL_BUG_REPORT=bugs@readersourcing.example
EMAIL_ADMIN=contact@readersourcing.example
PUBLIC_BASE_URL=https://readersourcing.example
CORS_ALLOWED_ORIGINS=https://client.readersourcing.example
FORCE_SSL=true
RAILS_LOG_TO_STDOUT=true
```

The value of `PUBLIC_BASE_URL` is an origin: it contains the scheme, host, and optional port, but no path, query string, fragment, or credentials. Production refuses to start when the value is missing or invalid. Its host is accepted by the Rails host authorization policy; further trusted hosts can be listed through `ADDITIONAL_ALLOWED_HOSTS`. The public origin

is used for password recovery and prepared publication links rather than deriving either value from an untrusted request host.

`CORS_ALLOWED_ORIGINS` contains exact origins separated by commas. Supplying `*` deliberately permits requests from every origin and is inappropriate for an ordinary public deployment. `RS_Rate` requests browser permission for the selected server origin and therefore does not depend on a generated extension origin being listed here. Set `FORCE_SSL=true` only after HTTPS is available at the public origin.

Publication downloads reject private and reserved network addresses by default. The `RS_PDF_ALLOW_PRIVATE_NETWORKS` override is intended only for a controlled deployment in which trusted publications are deliberately served from such a network.

Prepared publications are stored below `RS_PDF_STORAGE_ROOT` and not below `public`. The publication path fields describe the current reader's copy as logical compatibility metadata, while `pdf_download_url` and `pdf_download_url_link` contain signed URLs for its original and annotated files with the configured lifetime. Replicas that serve the same prepared copies must share this storage and support atomic file renaming and file locks. Logging out does not revoke an issued download URL, which remains usable only until that short lifetime ends. Existing files under `public/user` can be moved once with `bin/rails publications:migrate_private_storage`; the task stops rather than combining an existing destination with the old tree.

Rate limit counters use the configured Rails cache store. The supplied Puma configuration runs a single process and uses a store in memory. Deployments with multiple server processes or replicas must configure a shared Active Support cache store so that every instance contributes to the same request budgets.

### 3.5.5 Sending Mails

`RS.Server` supports standard SMTP services for account confirmation, password recovery, password change notifications, bug reports, and contact messages. Recovery messages contain the link described above, which can be used once; the new password is entered only in `RS.Server` and is not included in either recovery or confirmation mail. The SMTP provider must expose a host, port 587 with STARTTLS, a user name, a password, and a verified sender domain. API keys specific to a provider can be placed in `SMTP_PASSWORD` when the provider uses an API key as its SMTP credential; the provider documentation defines the required values.

### 3.5.6 Connecting To The Database

A full connection string supplied through `DATABASE_URL` takes precedence over the individual `POSTGRES_*` variables in development and production. The test environment follows the same rule using `TEST_DATABASE_URL` and `POSTGRES_TEST_DB`.

```
ENV['DATABASE_URL'] ||  
  "postgresql://#{ENV['POSTGRES_USER']} | 'postgres'}:" \  
  "#{ENV['POSTGRES_PASSWORD']}@" \
```

```
"#{ENV['POSTGRES_HOST']} || 'localhost'}/" \
"#{ENV['POSTGRES_DB']} || 'rs_server'"
```

### 3.5.7 Logging To The Standard Output

Development writes logs to standard output by default. In production, setting `RAILS_LOG_TO_STDOUT` configures a tagged logger on standard output, which is appropriate for Docker and other distributed deployments.

## 4 RS\_PDF

RS.PDF [6] is the software library used by RS.Server to edit PDF files and add the required URL when a reader requests to save a publication for later reading. The current implementation appends a final page containing a clickable rating URL and a QR code, and stores the same URL in the PDF `BaseUrl` metadata field. In this integration, RS.PDF treats the complete URL as opaque data: the paper rating reference is created and later validated only by RS.Server.

Its command line interface allows RS.Server to invoke it directly. Because the components are deployed together, they do not require a separate network service or communication protocol. The executable is distributed as a single shaded JAR named `RS_PDF-v2.0.0-shaded.jar` and bundled in the RS.Server `lib` directory.

### 4.1 Implementation and Technology

RS.PDF is developed in Kotlin, an object oriented programming language that inter-operates with the Java Virtual Machine. This interoperability allows developers to reuse software distributed in JAR format and call Java classes directly from Kotlin. The current build targets Java 21 and uses Kotlin 2.4.

This programming language was chosen because it incorporates many modern features and receives robust support. More importantly, the underlying tool used to edit PDF files is *PDFBox*,<sup>11</sup> a Java library that provides a complete PDF manipulation toolkit. Therefore, RS.PDF serves as a wrapper for PDFBox and adds the required reference to the PDFs requested by readers. ZXing<sup>12</sup> generates the QR code placed beside the link annotation, Apache Commons CLI parses the execution parameters, and Log4j 2 provides logging.

Kotlin was created by JetBrains,<sup>13</sup> which announced full support for Android development with Google in 2017.<sup>14</sup> In the same year, JetBrains also announced Kotlin/Native, which compiles Kotlin programs to native binaries without requiring the JVM.

---

<sup>11</sup><https://pdfbox.apache.org/>

<sup>12</sup><https://github.com/zxing/zxing>

<sup>13</sup><https://www.jetbrains.com/>

<sup>14</sup><https://blog.jetbrains.com/kotlin/2017/05/kotlin-on-android-now-official/>




Figure 6: Package diagram of RS\_PDF.

The Kotlin documentation includes an official comparison with Java,<sup>15</sup> and the language has been discussed extensively by software developers.<sup>16</sup>

## 4.2 Package Diagram

Figure 6 shows how RS\_PDF is divided into packages and provides a general overview of its internal architecture.

Interaction with RS\_Server begins in the **program** package. The *Program* object parses and validates the command line options before delegating to the publication controller.

The **utils** package provides shared filesystem defaults, program constants, and logging configuration. Server addresses are supplied through the required URL argument and are not embedded as constants in the code.

The **publications** package contains the PDF business logic. Its Controller/Model separation follows the same object oriented intent as MVC without depending on Rails: the controller discovers one or more PDF files and creates the corresponding models, while each model loads a document, adds the link and QR code, writes metadata, and saves a new file. A View is unnecessary because the result is the edited PDF itself.

## 4.3 Class Diagram

Figure 7 shows the main classes of RS\_PDF. The structure follows Controller/Model responsibilities: *Program* owns process parsing and flow, *PublicationController* manages the collection, and *PublicationModel* performs PDF transformations through PDFBox and ZXing.

<sup>15</sup><https://kotlinlang.org/docs/reference/comparison-to-java.html>

<sup>16</sup><https://medium.com/@octskyward/why-kotlin-is-my-next-programming-language-c25c001e26e3>




Figure 7: General class diagram of RS\_PDF v2.0.0.

The *Parameters* data class transports the caption, URL, optional authentication token, and optional publication identifier through the controller boundary. This keeps the *PublicationModel* interface stable as related input data evolve.

## 4.4 Installation

RS\_PDF is bundled with compatible RS\_Server distributions, eliminating the need for a separate installation. To use it independently, download the shaded JAR from the repository releases<sup>17</sup> or build it from source:

```
./mvnw clean verify
```

On Windows use `mvnw.cmd clean verify`. The Maven Wrapper selects the configured Maven version and the executable is generated as `target/RS_PDF-v2.0.0-shaded.jar`.

### 4.4.1 Requirements

- Java runtime  $\geq 21$  to execute the shaded JAR;
- JDK  $\geq 21$  to compile and verify the project from source.

### 4.4.2 Command Line Interface

The behavior of RS\_PDF is configured through the command line options in Table 3. Caption and URL are always required. Input and output paths may be omitted only together, in which case the default `data` and `res` directories are used. The authentication token and publication identifier are optional, but if either is supplied then both custom paths and both server values must be supplied.

<sup>17</sup>[https://github.com/Miccighel/Readersourcing-2.0-RS\\_PDF](https://github.com/Miccighel/Readersourcing-2.0-RS_PDF)

| Short              | Long                         | Description                                                              | Constraint                                       |
|--------------------|------------------------------|--------------------------------------------------------------------------|--------------------------------------------------|
| <code>-pIn</code>  | <code>--pathIn</code>        | Input PDF file or directory, using a relative or absolute path.          | Optional; requires <code>-pOut</code> .          |
| <code>-pOut</code> | <code>--pathOut</code>       | Existing output directory for annotated PDF files.                       | Optional; requires <code>-pIn</code> .           |
| <code>-c</code>    | <code>--caption</code>       | Caption of the clickable link added to the PDF.                          | Required.                                        |
| <code>-u</code>    | <code>--url</code>           | Rating URL used by the link, QR code, and <code>BaseUrl</code> metadata. | Required valid URL.                              |
| <code>-a</code>    | <code>--authToken</code>     | Authentication token received from <code>RS_Server</code> .              | Optional; requires paths and <code>-pId</code> . |
| <code>-pId</code>  | <code>--publicationId</code> | Publication identifier received from <code>RS_Server</code> .            | Optional; requires paths and <code>-a</code> .   |

Table 3: Command line options of RS\_PDF.

The following example assumes that:

- there is a folder containing  $n$  PDF files at `C:\data`;
- annotated files must be saved in the existing folder `C:\out`;
- the shaded JAR is located at `C:\lib\RS_PDF-v2.0.0-shaded.jar`.

RS\_PDF can then be executed with the following command:

```
java -jar C:\lib\RS_PDF-v2.0.0-shaded.jar \
-pIn C:\data -pOut C:\out \
-c "Express your rating" \
-u "https://readersourcing.example/rate"
```

## 5 RS\_Rate

RS\_Rate [8] is an extension designed to function as a client for the Readersourcing 2.0 ecosystem without requiring access to its website. Compatible with Google Chrome,<sup>18</sup> Microsoft Edge,<sup>19</sup> and Mozilla Firefox,<sup>20</sup> the extension allows users to rate publications directly from their browsers. This eliminates the need to navigate to the main website, streamlining the process of providing ratings for publications.

<sup>18</sup><https://www.google.com/chrome/>

<sup>19</sup><https://www.microsoft.com/en-us/edge/>

<sup>20</sup><https://www.mozilla.org/firefox/>

The primary objective of RS\_Rate is to enable readers to rate a publication with minimal effort: just a few clicks or keystrokes within the Readersourcing 2.0 ecosystem. In doing so, it contributes to a more dynamic online reading experience. RS\_Rate is the first client developed for the project beyond the web interface available on the main portal.

Looking ahead, our vision includes expanding the compatibility of RS\_Rate to Safari and other popular browsers. Our commitment is to make this rating tool accessible across a broad range of browsers, ensuring users can seamlessly interact with content and provide feedback regardless of their preferred browser. The current Chromium and Firefox targets share the same application source and rating logic.

## 5.1 Implementation and Technology

RS\_Rate is an extension for Google Chrome, Microsoft Edge, and Mozilla Firefox. These extensions are developed using standard web technologies such as HTML, CSS, and JavaScript. They are collections of files packaged according to the requirements of each browser. The build process copies and checks all executable assets locally, then produces unpacked Manifest V3 extensions in `dist/chromium` and `dist/firefox`.

The JavaScript component also uses jQuery to simplify DOM selection and manipulation, event management, animations, and AJAX operations. RS\_Rate relies on these AJAX features to provide responsive interactions with the server. Bootstrap 4 and its plugins based on jQuery provide the markup and interaction behavior used by the extension views. Dependencies remain declared in `package.json` and are resolved through the committed Yarn lockfile.

## 5.2 Installation

The original 2019 release of RS\_Rate remains available on the Google Chrome Web Store. The current development source produces the Manifest V3 packages described in the following section. To install the original release with a browser based on Chromium, use the store page:

- **Google Chrome and Microsoft Edge:** <https://chromewebstore.google.com/detail/readersourcing-20-rsrate/hlkdlnpgijhdkbdlhmgeemffaocjagg>

The repository also produces a Firefox Manifest V3 package ready for local validation and submission to Mozilla Add-ons. It uses the following stable Gecko identifier:

`rs-rate@readersourcing.org`

For local testing, open `about:debugging`, choose *This Firefox*, and load

`dist/firefox/manifest.json`

using *Load Temporary Add-on*. Distribution through Mozilla Add-ons requires validation and signing by Mozilla.

### 5.3 Development and Packaging

The development toolchain uses Node.js 24 LTS and Yarn 4.17.1. From the repository root:

```
corepack enable
yarn install --immutable
yarn verify
```

`yarn verify` builds both targets, checks that all runtime assets are local, validates the Firefox build with Mozilla’s `web-ext`, and runs the client contract tests. The Chromium and Firefox archives are generated with:

```
yarn package
```

The files `rs_rate-chromium.zip` and `rs_rate-firefox.zip` are written to `artifacts`. The commands `yarn package:chromium` and `yarn package:firefox` create the individual archives. The Firefox target requires Firefox 142 or later. Its manifest declares the account, browsing, website interaction, and publication content categories required by the Readersourcing workflow. The extension does not include an analytics channel.

Development builds initially target `http://localhost:3000`. Users can change the `RS_Server` host in the extension options, and the same configured host is used for normal API requests and PDF extraction. When the value is saved, the browser asks the reader to grant access to that HTTP or HTTPS origin. The permission covers only the selected origin, while any configured application path remains part of the server address.

### 5.4 Usage

Figure 8 shows a Google Chrome window with the extension active for a publication. This represents a typical reader visiting a publisher’s website to access a paper in PDF format. The figure also shows the first page displayed when the reader opens the client. From this page, the reader can reach either the login page shown in Figure 9 or the registration page. The login page also provides access to a password recovery page, which requests the email containing the link to the `RS_Server` password form.

If a reader has yet to register for Readersourcing 2.0, they can navigate from the page shown in Figure 8 to the one displayed in Figure 11 and fill in the registration form. Once they complete the standard registration and login operations, they will find themselves on the rating page, as shown in Figure 10.

In the central section of the rating page, a reader can use the slider to choose a value in the 0–100 interval. After selecting the desired value, they submit the rating by clicking the green *Rate* button. The options page also allows the reader to make the rating anonymous. Readers must still be logged in when submitting an anonymous rating so that the system can prevent spam and repeated ratings of the same publication. Apart from these checks, the rating is processed without exposing the reader’s identity.




Figure 8: RS\_Rate characterized as an extension having a popup action.

The image shows the login page of RS\_Rate. It has a blue header bar with a left arrow and a gear icon. The title "READERSOURCING 2.0" is centered, followed by the instruction "Insert your login data". There are two input fields: one for "Email (e.g. address@example.com)" and one for "Password". Below the password field, there are two links: "Did you forget your password?" and "Do you want to create a new account?". At the bottom, there is a green "LOGIN" button with a right-pointing arrow icon.

Figure 9: The login page of RS\_Rate.




Figure 10: The rating page of RS\_Rate.

Figure 11: The user registration page of RS\_Rate.




Figure 12: The rating page of RS\_Rate after a “save for later” request.

If the reader prefers to rate the publication later, they can use the *Download* button in the *Save For Later* section. Before the operation begins, the extension reports that the publication is ready to prepare. During the request, the same status area reports validation, annotation, and final verification. Once the procedure is complete, the prepared publication can be opened, as shown in Figure 12.

If the publication server requires access through the reader’s browser session or cannot be reached, the interface explains the condition and displays the local PDF alternative. The reader can upload the original file already present on the device without changing the publication URL used by the domain model. This source follows the same validation and verification procedure as a direct retrieval. Other conditions, including size limits, encrypted or malformed PDFs, annotation failures, and verification failures, are reported separately instead of being presented as a generic download error.

The reader can then download the annotated publication. The refresh button, located to the right of the *Download* button, retrieves a new copy of the publication, annotates it, and makes it available again. This feature is useful when an author updates a publication after the first download.

The downloaded PDF contains a new final page with the rating URL. The URL includes the encrypted paper rating reference associated with the reader and publication, not the reader’s JWT. The reference has no automatic expiration and remains independent of login and logout operations. Figure 13 shows an example opened in a PDF reader.

Following the reference takes the reader to a dedicated page in RS\_Server. This interface




Figure 13: A publication annotated through RS\_Rate.

allows them to rate the publication independently of the browser extension used to store the reference. For example, a reader can transfer the annotated publication to a tablet and use its integrated browser to submit the rating. Figure 14 shows the interface displayed after the reader follows the stored reference. The reader must have a current authenticated session, so another person who receives a copy of the publication cannot use the reference on their behalf. If the reader has logged out, they sign in again and reopen the same reference; the PDF itself does not need to be annotated again.

Whenever a reader rates a publication, RS\_Server updates the scores computed by both the RSM and TRM models. RS\_Rate displays the two publication scores at the bottom of the rating page, as shown in Figures 10 and 12. The reader can view their own scores by clicking the profile button in the upper right corner. The profile page shown in Figure 15 also allows the reader to change their password.

## 6 RS\_Py

RS\_Py [7] is an additional component of the Readersourcing 2.0 ecosystem, providing a fully functional Python implementation of the models presented by Soprano et al. [5]. These models are also encapsulated by the server application of Readersourcing 2.0, as described in Section 3.

Developers familiar with Python can use RS\_Py to generate and test new simulations of readers rating a set of publications. They can alter the internal logic of the models to test new approaches without having to fork and edit the full Readersourcing 2.0 implementation.

**READERSOURCING 2.0**

Rate this publication

**50**

☐ Anonymize

Validate yourself

Email (e.g. address@example.com)

johndoe@mail.com

Password

.....

**RATE**

Figure 14: The server interface to rate a publication.

**USER PROFILE** X

|                    |                  |
|--------------------|------------------|
| <b>First Name</b>  | John             |
| <b>Last Name</b>   | Doe              |
| <b>Email</b>       | johndoe@mail.com |
| <b>Score (RSM)</b> | 99.89/100        |

Figure 15: The profile page of RS\_Rate.

## 6.1 Implementation and Technology

RS.Py is a collection of *Jupyter Notebooks*. Jupyter<sup>21</sup> is an open source web application powered by Python<sup>22</sup> for creating and sharing documents containing live code, equations, visualizations, and narrative text. Notebooks are composed of cells that can be run independently, providing a detailed overview of the implemented computations.

## 6.2 Installation

To use the RS.Py notebooks, clone the repository<sup>23</sup> and place it somewhere on the filesystem. The project targets Python 3.13, and its tested direct dependencies are declared in `requirements.txt`. An isolated environment can be created as follows:

```
python -m venv .venv
source .venv/bin/activate
python -m pip install -r requirements.txt
```

On Windows, activate the environment with `.venv\Scripts\activate`.

## 6.3 Usage

RS.Py is organized into the following main folders:

- The `data` folder stores the datasets used to test the models presented by Soprano et al. [5] and the compressed synthetic datasets.
- The `models` folder stores the output of these models and the computed quantities.
- The `notebooks` folder contains Jupyter notebooks used to generate datasets, implement the models, and analyze experiments.
- The `scripts` folder is generated on demand and contains implementations of the Jupyter notebooks as pure Python scripts.
- The `src` folder contains the conversion utility and the domain implementation held in memory.
- The `tests` folder contains the numerical and utility characterization tests.
- The `utils` folder contains auxiliary tooling, including the power law generator.

Within the `notebooks` folder, the principal files are:

- `1.Seed-Power-Law.ipynb` generates synthetic data using power law distributions;

---

<sup>21</sup><https://jupyter.org/>

<sup>22</sup><https://www.python.org/>

<sup>23</sup><https://github.com/Miccighel/Readersourcing-2.0-RS-Py>

- `2.Readersourcing.ipynb` implements and runs the RSM model;
- `3.TrueReview.ipynb` implements and runs the TRM model;
- `4.Experiments.ipynb` supports experimental analysis;
- `ReadersourcingToolkit.py` contains shared Python functionality used by the notebooks.

The `src/Convert.py` script converts the numbered notebooks into pure Python scripts and stores them in the generated `scripts` folder together with their shared toolkit.

The behavior of the notebooks can be customized by modifying the parameter settings found in their initial cells. To run them, navigate to the `notebooks` folder and type `jupyter lab`. This command starts the Jupyter server and opens the interface from which the notebooks can be executed in numerical order.

The implementation held in memory under `src/rs_py` and its characterization tests reproduce the reader, publication, and rating behavior currently shared with RS\_Server. They complement rather than restrict the principal purpose of RS\_Py: developers remain free to alter the internal logic of the notebooks and test new approaches. Additional object oriented simulation work, including an explicit author entity, is maintained in the Readersourcing\_OO repository.<sup>24</sup>

## 7 Overview

Figure 16 summarizes the capabilities of Readersourcing 2.0. In this example, four readers—RD1, RD2, RD3, and RD4—use RS\_Rate to rate publication P1. Readers RD1, RD2, and RD3 choose *Save for later* and receive a version of the publication annotated with a link, denoted by P1+Link.

After some time, RD1 opens P1+Link with a preferred PDF reader. RD2 sends it to a tablet, while RD3 opens it in a desktop browser. By following the link or scanning the QR code added through RS\_PDF, they reach the dedicated RS\_Server page and provide their ratings. In contrast, RD4 rates P1 immediately through the RS\_Rate interface.

---

<sup>24</sup>[https://github.com/EddyMaddalena/Readersourcing\\_OO](https://github.com/EddyMaddalena/Readersourcing_OO)




Figure 16: Readers' interaction modalities with the Readersourcing 2.0 ecosystem.

## References

- [1] Luca De Alfaro and Marco Faella. TrueReview: A Platform for Post-Publication Peer Review. In: *CoRR* abs/1608.07878 (2016). arXiv: 1608.07878. URL: <http://arxiv.org/abs/1608.07878>.
- [2] Martin Fowler. *UML Distilled: A Brief Guide to the Standard Object Modeling Language*. 3rd ed. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc., 2003. ISBN: 0321193687.
- [3] Erich Gamma et al. *Design Patterns: Elementi per il Riutilizzo di Software ad Oggetti*. Pearson, Jan. 2002.
- [4] Stefano Mizzaro. Quality control in scholarly publishing: A new proposal. In: *Journal of the American Society for Information Science and Technology* 54.11 (2003), pp. 989–1005. DOI: 10.1002/asi.10296. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/asi.10296>.
- [5] Michael Soprano and Stefano Mizzaro. Crowdsourcing Peer Review: As We May Do. In: *Digital Libraries: Supporting Open Science*. Vol. 988. Communications in Computer and Information Science. Springer, 2019, pp. 259–273. DOI: 10.1007/978-3-030-11226-4\_21. URL: [https://doi.org/10.1007/978-3-030-11226-4\\_21](https://doi.org/10.1007/978-3-030-11226-4_21).
- [6] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS\_PDF*. DOI: 10.5281/zenodo.1442597. URL: <https://doi.org/10.5281/zenodo.1442597>.

- [7] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS\_Py*. DOI: 10.5281/zenodo.3245208. URL: <https://doi.org/10.5281/zenodo.3245208>.
- [8] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS\_Rate*. DOI: 10.5281/zenodo.1442599. URL: <https://doi.org/10.5281/zenodo.1442599>.
- [9] Michael Soprano and Stefano Mizzaro. *Readersourcing 2.0: RS\_Server*. DOI: 10.5281/zenodo.1442630. URL: <https://doi.org/10.5281/zenodo.1442630>.
- [10] Michael Soprano, Kevin Roitero, and Stefano Mizzaro. HITS Hits Readersourcing: Validating Peer Review Alternatives Using Network Analysis. In: *Proceedings of the 4th Joint Workshop on Bibliometric-enhanced Information Retrieval and Natural Language Processing for Digital Libraries*. Vol. 2414. CEUR Workshop Proceedings. 2019, pp. 70–82. URL: <https://ceur-ws.org/Vol-2414/paper7.pdf>.
- [11] Michael Soprano et al. Evaluation of Crowdsourced Peer Review using Synthetic Data and Simulations. In: *Proceedings of the 21st Conference on Information and Research Science Connecting to Digital and Library Science*. Vol. 3937. CEUR Workshop Proceedings. 2025. URL: <https://ceur-ws.org/Vol-3937/paper8.pdf>.